

Indirect Dynamic Mechanism Design with Endogenous Implementation Noise: Identifying Rational Types under Observational Pooling^{*}

Humberto Bernal

*Economics Program, Universidad Colegio Mayor de Cundinamarca; Ph.D. in Economics,
Universidad de los Andes, Colombia.*

This paper studies dynamic indirect mechanism design when a principal must screen an agent whose equilibrium actions pool, so passive observation cannot separate types. The contribution is theoretical. The paper introduces the Mano de Dios–Guadalupe (MDG) process, an endogenous-noise implementation kernel that maps latent strategic intentions into executed actions while preserving full support off the equilibrium path without an exogenous trembling-hand perturbation. Combining MDG with competing risks and an active information-acquisition motive, the paper establishes a Type Identification Theorem: uniform total-variation separation between type-contingent observation laws implies mutual singularity of trajectory laws by the criterion of [Kakutani \(1948\)](#), so Bayesian beliefs converge almost surely to the true type on the prolonged-interaction event. A Conditional Mechanism Implementability Theorem then guarantees a perfect Bayesian equilibrium satisfying incentive compatibility and individual rationality whenever the feasible policy set is non-empty. Both results are machine-checked in Lean 4. Kidnapping for ransom in Colombia is used as an applied validation: microeconomic fingerprints calibrate type-specific hazards, and ELN–PAR simulations illustrate finite-horizon learning, instrument paths, and constraint compliance under the calibrated mechanism.

Keywords: indirect mechanism design, dynamic screening, endogenous implementation noise, Bayesian learning, type identification

JEL classification: C73, D82, D86

^{*}Economist, professor of the economics program at the Universidad Colegio Mayor de Cundinamarca and Ph.D. in economics from the Universidad de los Andes, Colombia. E-mail: hbernal@uniandes.edu.co, hbernal@universidadmayor.edu.co. ORCID: 0000-0002-4864-1726. All research is the sole responsibility of the author, received no funding, and was developed as an independent initiative.

1. Introduction and Related Literature

Mechanism design asks how a principal can implement an allocation when agents hold private information and act strategically. In the classical program of [Gibbard \(1973\)](#) and [Myerson \(1979, 1981\)](#), the Revelation Principle represents any incentive-compatible outcome of an indirect mechanism by an equivalent direct mechanism in which truthful reporting is optimal. That representation is a theoretical reduction, not an institutional prescription: when messages are restricted, actions are noisy, or types refuse literal revelation, the economically relevant object remains an *indirect* mechanism based on observable actions, public histories, and implementable instruments ([Jackson, 2001](#); [Bergemann and Morris, 2012](#)). The present paper develops such an object for dynamic environments in which equilibrium behavior pools: distinct types generate observationally similar short-run signals, so a myopic policy that minimizes current loss can freeze the very pooling that prevents identification tomorrow.

The central design problem is therefore threefold. First, how should the principal specify an action-based indirect mechanism when the message space is not a free report of type? Second, how can implementation frictions be modeled endogenously so that Bayes' rule remains well defined off path, without relying on an exogenous trembling-hand device? Third, under what conditions do public signals identify the agent's true type, and when is the resulting policy rule implementable in perfect Bayesian equilibrium (PBE)? The paper is situated at the frontier of dynamic mechanism design, expanding the foundational theoretical frameworks of [Courty and Li \(2000\)](#), [Eső and Szentes \(2007\)](#), and [Pavan et al. \(2014\)](#) by incorporating active informational exploration and structural noise, which are indispensable for its viability in adversarial institutional contexts.

1.1. Contribution relative to the literature on indirect mechanisms

Relative to the static implementability frontier defined by [Myerson and Satterthwaite \(1983\)](#), [Rochet \(1987\)](#), and [Rochet and Choné \(1998\)](#), and complementing the general existence conditions established by [Kadan et al. \(2017\)](#), the central contribution of this work lies in formulating and characterizing an indirect dynamic mechanism subject to endogenous implementation noise. This approach introduces an intrinsic stochastic friction between the agent’s latent strategic intention and the publicly recorded action, enabling both the Bayesian identification of rational types and conditional implementability under the perfect Bayesian equilibrium (PBE) criterion. Specifically, the paper structures this contribution through four interconnected theoretical results.

Indirect action-based dynamic mechanism.. The paper formalizes a dynamic indirect mechanism in which the principal chooses instruments as functions of public history and beliefs, while private agents best-respond in their information sets. Types need not be reported; they are inferred from the law of executed actions and outcomes. This preserves the Revelation Principle as a representation theorem while insisting that, for the environments considered here, the operational design object is the indirect rule.

MDG endogenous implementation kernel.. The Mano de Dios–Guadalupe (MDG) process—equivalently, an endogenous implementation kernel—maps latent intentions a_t^{i*} into executed actions $\tilde{a}_t^i$ through a hybrid-temperature logit with full support. Unlike classical refinements that inject exogenous trembles into strategies ([Selten, 1975](#); [Kreps and Wilson, 1982](#)), MDG treats implementation noise as an intrinsic friction of materialization. Under a positive residual floor, the ker-

nel maintains absolute continuity in the sense of [Kalai and Lehrer \(1993\)](#), keeping Bayesian updating well defined on every branch while concentrating mass on the intended action. A maximum-entropy characterization in the spirit of [Jaynes \(1957\)](#) disciplines the functional form.

Type Identification Theorem for rational types.. When each type induces a structural fingerprint that separates daily observation laws uniformly in total variation, the product-measure criterion of [Kakutani \(1948\)](#) yields mutual singularity of trajectory laws. Combined with the martingale stability of beliefs ([Doob, 1953](#); [Blackwell and Dubins, 1962](#)), posterior mass concentrates almost surely on the true type on the prolonged-interaction event $\mathcal{T} = \{\tau = \infty\}$. The theorem converts a local daily separation condition into global identification. Its limitation is explicit: asymptotic consistency is stated on $\mathcal{T}$; finite-horizon simulations in the application illustrate directional learning without substituting for the asymptotic result.

Conditional PBE implementability.. Identification alone does not make a mechanism executable. The Conditional Mechanism Implementability Theorem shows that, whenever the feasible set defined by incentive compatibility and individual rationality remains non-empty, there exists a PBE in which the principal’s optimal rule implements the social choice subject to those constraints. Thus MDG, learning, and incentive feasibility form a single architecture: noise enables off-path updating; separation enables identification; and non-empty $\Gamma_t(\mu_t)$ enables equilibrium implementation.

Applied validation.. Sections [4–5](#) map the abstract environment to kidnapping for ransom in Colombia. Multinomial logit and competing-risks estimates from

National Center for Historical Memory microdata supply type-specific hazards; ELN and PAR calibrations then report posterior paths, instrument trajectories, informational gain, and IR/IC compliance. The application is a validation layer, not the theoretical primitive.

The remainder of the article proceeds as follows. Section 2 presents the general indirect mechanism, the MDG kernel, competing risks, and best responses for arbitrary unknown types. Section 2.6 states the Type Identification Theorem. Section 3 establishes conditional PBE implementability and policy implications. Section 4 specializes the environment to Colombian kidnapping and anchors parameters empirically. Section 5 reports calibration and simulations. Section 6 concludes. Supplemental Appendices contain macroeconometrics, microeconomic detail, and Lean 4 proofs.

2. General Environment and Indirect Mechanism

This section develops the theoretical model in general form. A principal designs an indirect dynamic mechanism when an agent's private type is unknown and equilibrium actions may pool. Types need not be elicited by direct reports: the mechanism operates on public histories, executed actions, and instruments. The Revelation Principle of Gibbard (1973) and Myerson (1979, 1981) continues to imply that any incentive-compatible indirect outcome admits a direct representation; the contribution here is to construct and analyze the indirect object itself when messages are restricted to materialized actions subject to endogenous implementation noise. Kidnapping for ransom appears only later (Section 4) as a concrete specialization of the same primitives.

Primitives (general notation). Let $\mathcal{J}$ be a finite set of players containing a principal P and privately informed agents. Each episode is indexed by discrete time $t \in \mathbb{T}$ until an endogenous stopping time τ . A type $\theta \in \Theta$ is drawn from a common prior $\mu_0 \in \Delta(\Theta)$ with full support; only a designated agent observes the strategically relevant component of θ . In each period the principal chooses instruments x_t as a measurable function of public history h_t and beliefs μ_t ; agents choose latent intentions a_t^{i*} ; an implementation kernel $Q_t(\tilde{a}_t^i \mid a_t^{i*}, x_t, h_t)$ —the MDG process—delivers executed actions $\tilde{a}_t^i$; public signals ζ_t update μ_{t+1} by Bayes’ rule. The design goal is to implement a social choice rule that is sequentially rational, incentive compatible, and individually rational while inducing informative laws of observation.

Specialization used below. For concreteness and continuity with the Lean-checked theorems, the remainder of this section writes the general objects in the language of a three-player adversarial screening game $\{S, K, F\}$ (principal, informed adversary, and complementary agent), with private type $\theta_K \in \Theta_K$, instruments (α_t, γ_t) , and competing-risk outcomes. Nothing in the identification or implementability arguments requires the labels “state,” “captor,” or “family”: those names are interpretive. The Colombian type menu $\{\text{DC}, \text{PAR}, \text{ELN}, \text{FARC}\}$ is introduced only when the application maps Θ_K to estimated criminal technologies.

2.1. Definition of the mechanism

Let $\mathcal{J} = \{S, K, F\}$ be the set of players in the specialized notation, where S is the principal (state), K the privately informed agent (adversary), and F a complementary agent (family). Each episode is characterized by the vector of

types $\theta = (\theta_K, \theta_F, \theta_V, \theta_S) \in \Theta$, where the strategically private component θ_K takes values in a finite menu Θ_K . In the Colombian application of Section 4, $\Theta_K = \{\text{DC}, \text{PAR}, \text{ELN}, \text{FARC}\}$ indexes distinct criminal technologies—common delinquency, paramilitary groups, National Liberation Army, and Revolutionary Armed Forces of Colombia—but the theorems treat Θ_K as an arbitrary finite type space; $\Theta_F = \{\text{High}, \text{Low}\}$ represents the family’s payment capacity; $\Theta_V = \{\text{Public}, \text{Private}\}$ represents the victim’s profile; and $\Theta_S = \{\text{Hard}, \text{Lax}\}$ represents the institutional environment. Each episode is also localized in an observable geographic zone $z \in \mathcal{R}$, where $\mathcal{R}$ groups five regions (Metropolitan, Andean, Caribbean, Pacific, and East-Jungle); z is not a private strategic type but a public initial condition that disciplines the prior over the captor’s type. Only θ_K is private information of K ; the other components and z are observable upon opening the episode, with an initial belief $\mu_0(\theta \mid z, \theta_V) = \mathbb{P}(\theta_K = \theta \mid z, \theta_V) \in \Delta(\Theta_K)$.

The public history $h_t \in H_t$ recursively records, up to the beginning of t and without including the type vector θ or the belief μ_t , the following elements: t is the period index; $\alpha_{[0,t-1]}^*$ and $\gamma_{[0,t-1]}^*$ are the trajectories of the state’s financial interdiction and operational pressure instruments; $\tilde{a}_{[0,t-1]}^F$, $\tilde{a}_{[0,t-1]}^K$, and $\tilde{a}_{[0,t-1]}^S$ are the actions executed by the family, the captor, and the state; $m_{[0,t-1]}$ is the sequence of physical outcomes; and $d_{[0,t-1]}$ is the sequence of institutional detection signals, where $d_t \in \{0, 1\}$ indicates whether in period t the state detected collusion between the family and the captor, i.e., a clandestine ransom payment bypassing the financial interdiction policy.

$$h_t = (t, \alpha_{[0,t-1]}^*, \gamma_{[0,t-1]}^*, \tilde{a}_{[0,t-1]}^F, \tilde{a}_{[0,t-1]}^K, \tilde{a}_{[0,t-1]}^S, m_{[0,t-1]}, d_{[0,t-1]}), \quad (1)$$

$$h_{t+1} = (h_t, \alpha_t^*, \gamma_t^*, \tilde{a}_t^F, \tilde{a}_t^K, \tilde{a}_t^S, m_t, d_t). \quad (2)$$

The information sets when choosing in t are

$$\begin{aligned}\mathcal{I}_t^S &= (h_t, z, \theta_F, \theta_V, \theta_S, \mu_t), \\ \mathcal{I}_t^F &= (h_t, z, \theta_F, \theta_V, \theta_S), \quad \mathcal{I}_t^K = (h_t, z, \theta_K, \theta_F, \theta_V, \theta_S).\end{aligned}$$

Although μ_t is internal to the state, its exclusion from the private information sets does not generate inconsistency: in the PBE, private agents share the same prior and the same public history h_t , and therefore replicate the Bayesian updating of the planner [Fudenberg and Tirole \(1991a\)](#).

The space of institutional allocations of the state is $\mathcal{X}_t = \mathcal{A}^{S,\text{disc}} \times \mathcal{A}^\alpha \times \mathcal{A}^\gamma$. In each period t , players formulate optimal strategic intentions a_t^{i*} ; the MDG process transforms them into executed actions $\tilde{A}_t = (\tilde{a}_t^F, \tilde{a}_t^K, \tilde{a}_t^S)$ that enter the public registry and determine h_{t+1} and the public signal ζ_t . Conditional on the executed actions and the optimal institutional policy, the system determines the physical outcome $m_t \in \mathcal{M}_{\text{out}}$ through a block of competing risks, with $\mathcal{M}_{\text{cont}} = \{\text{cont}\}$, where cont denotes the continuation of captivity, $\mathcal{M}_{\text{term}} = \{\text{pay}, \text{res}, \text{rel}, \text{kill}\}$, which groups the termination causes: ransom payment, tactical rescue, voluntary release, and death of the victim, respectively, and $\mathcal{M}_{\text{out}} = \mathcal{M}_{\text{cont}} \cup \mathcal{M}_{\text{term}}$.

The horizon of the episode is not fixed in advance but is determined by the captivity dynamics themselves: the episode concludes in the first period in which a terminal outcome materializes, so the effective horizon is the endogenous stopping time

$$\tau = \inf\{t \geq 0 : m_t \in \mathcal{M}_{\text{term}}\}, \quad (3)$$

where the recursive construction of h_t ensures that the event $\{\tau = t\}$ is measurable with respect to the public history at the beginning of t , guaranteeing that the state recognizes the closure of the episode in real time. On these elements, with $\mathbb{T} =$

$\{0, 1, \dots, \tau\}$ as the active periods index, $\Theta = \Theta_K \times \Theta_F \times \Theta_V \times \Theta_S$ the type space, $(H_t)_{t \in \mathbb{T}}$ the family of public history spaces recursively generated by (1)–(2), $(\mu_t)_{t \in \mathbb{T}}$ the belief process in $\Delta(\Theta_K)$ updated by Bayes' rule throughout the episode, and $\mathbb{P}$ the ex-ante law induced by the strategy profile s , the prior μ_0 , and the captivity dynamics, the dynamic indirect mechanism is formally defined as the tuple

$$\mathcal{D} = \langle \mathcal{J}, \mathbb{T}, \Theta, \mathcal{R}, (\mathcal{X}_t)_{t \in \mathbb{T}}, (H_t)_{t \in \mathbb{T}}, (\mu_t)_{t \in \mathbb{T}}, \mathbb{P} \rangle.$$

Given $z \in \mathcal{R}$ and $\theta_K \sim \mu_0$, the principal's strategy $s^S = \chi = (\chi_t)_{t \in \mathbb{T}}$ is implemented via the measurable public policy rule $\chi_t: H_t \times \mathcal{R} \rightarrow \mathcal{A}^{S, \text{disc}} \times \mathcal{A}^\alpha \times \mathcal{A}^\gamma$, which interacts with the private responses of K and F in their information sets; the social choice rule is formalized by the measurable mapping $f_t: \mathcal{I}_t^S \rightarrow \mathcal{X}_t$, such that the policy vector in equilibrium is $x_t^* := f_t(\mathcal{I}_t^S) = (a_t^{S*}, \alpha_t^*, \gamma_t^*) \in \mathcal{X}_t$. This rule is implementable if it constitutes the path of a PBE in the sense of [Fudenberg and Tirole \(1991a\)](#), which guarantees that the informed agent cannot systematically induce the principal to calibrate instruments on a type that is not the true one; implementability is achieved when the pair (s, μ) , with $s = (s^S, s^K, s^F)$ and $\mu_t \in \Delta(\Theta_K)$, satisfies two conditions in every reachable information set: *sequential rationality*, whereby each agent $i \in \{K, F\}$ maximises $\mathcal{U}_t^i$ given $\mathcal{I}_t^i$ and the strategies of the others while the state minimizes the expected social loss weighted by μ_t ; and *Bayesian consistency*, whereby in every node reached with positive probability under $\mathbb{P}^s$, the beliefs satisfy

$$\mu_t(\theta) = \mathbb{P}^s(\theta_K = \theta \mid \mathcal{F}_t^{\text{pub}}), \quad \mathcal{F}_t^{\text{pub}} = \sigma(z, \zeta_0, \dots, \zeta_{t-1}),$$

where $\mathcal{F}_t^{\text{pub}}$ is the public filtration at the beginning of t , generated by the σ -algebra of the initial geographic location z and the public signals $\zeta_0, \dots, \zeta_{t-1}$; each signal $\zeta_u = (\alpha_u^*, \gamma_u^*, \tilde{a}_u^F, \tilde{a}_u^K, \tilde{a}_u^S, m_u, d_u)$ aggregates all observables of period u and

constitutes the informational support over which the state updates its beliefs about θ_K ; so that the instruments (α_t^*, γ_t^*) are calibrated on correct inferences about the captor's type.

2.2. Summary of the mechanism

In each active period t , the protocol advances in three sequential stages. In the first, the agents formulate their optimal intentions: the state resolves $x_t^* = f_t(\mathcal{I}_t^S)$, and private agents select their best responses in their respective information sets, conditioned on the observable public policy (α_t^*, γ_t^*) signaled by the state. In the second, the MDG process transforms these intentions into executed actions $\tilde{A}_t$, guaranteeing full support across all branches of the game. In the third, the material state of the episode

$$\mathcal{C}_t(\theta_K) = (\theta_K, \theta_F, \theta_V, \theta_S, z, \tilde{a}_t^F, \tilde{a}_t^K, \tilde{a}_t^S, \alpha_t^*, \gamma_t^*, p_{\text{det},t}),$$

with $p_{\text{det},t}$ as the logit probability of detecting collusion, which is increasing in the state's instruments. This material state feeds a block of competing risks that determines m_t and d_t ; the public signal ζ_t aggregates all observables of the period and updates the history to h_{t+1} and the belief to μ_{t+1} through Bayes' rule.

2.3. Probabilistic technology

This section formalizes the stochastic architecture that transforms strategic plans into observable realizations. In the base specification, the MDG process takes as input the latent strategic intention a_t^{i*} and generates an executed action $\tilde{a}_t^i$ through an implementation law concentrated on the plan, but with full support over $\mathcal{A}^i$. The executed profile $\tilde{A}_t = (\tilde{a}_t^K, \tilde{a}_t^S, \tilde{a}_t^F)$ is the observable variable that

enters the public record, while the strategic intention remains latent. Given $a_t^{i*} \in \mathcal{A}^i$ and the state vector X_t , execution is determined by

$$\tilde{a}_t^i \sim \mathbb{P}_{I,i}(\tilde{a}_t^i \mid a_t^{i*}, X_t),$$

and, for each agent $i \in \{K, S, F\}$ and action $a \in \mathcal{A}^i$, it is specified as

$$\mathbb{P}_{I,i}(\tilde{a}_t^i = a \mid a_t^{i*}, X_t) = \frac{\exp(\mathbf{1}\{a = a_t^{i*}\}/T_t)}{\sum_{a' \in \mathcal{A}^i} \exp(\mathbf{1}\{a' = a_t^{i*}\}/T_t)}, \quad (4)$$

where $T_t > 0$ in the base specification is the temperature of the system: it measures the implementation friction of the episode, not the agent's uncertainty about its own type. To capture strategic learning and institutional maturation jointly, define

$$T_t = T_0 \max\left\{\frac{H(\mu_t)}{H(\mu_0)} e^{-\eta_{\text{cal}} t}, \underline{c}\right\}, \quad (5)$$

with $T_0 > 0$ as the initial temperature, $\eta_{\text{cal}} > 0$ the calendar decay rate, $H(\mu_t) = -\sum_{\theta} \mu_t(\theta) \ln \mu_t(\theta)$ the entropy of [Shannon \(1948\)](#), and $\underline{c} \in [0, 1]$ the floor of structural noise. In the base specification, the logit distribution (4) satisfies dominance, full support, and maximum entropy of [Shannon \(1948\)](#) subject to those constraints; under the characterization theorem of [Jaynes \(1957\)](#), this functional form is the entropic maximizer. Dominance concentrates most of the mass on a_t^{i*} , while every alternative action retains strictly positive probability as implementation friction. Thus, the MDG fulfills a function equivalent, on the operational plane, to the trembling hand of [Selten \(1975\)](#) and the sequential evaluation system of [Kreps and Wilson \(1982\)](#): no branch of the game is excluded by construction, and beliefs can be updated by Bayes at reached information sets, in line with the consistency of [Fudenberg and Tirole \(1991b\)](#).

Bounded institutional maturation makes the informational component of the noise decline as the system accumulates information, although the structural floor

$\underline{c}$ persists. When $\underline{c} > 0$, even if $H(\mu_t) \rightarrow 0$ the probability of executing exactly a_t^{i*} remains strictly less than one: the floor preserves full support over $\mathcal{A}^i$ and no branch is excluded in the limit. The special case $\underline{c} = 0$ is interpreted as a limit without residual noise: it allows full convergence $\mathbb{P}(\tilde{a}_t^i = a_t^{i*} \mid \mathcal{F}_t^{\text{pub}}) \rightarrow 1$ when $H(\mu_t) \rightarrow 0$, but sacrifices that full support. The base version with $\underline{c} > 0$ aligns with the absolute continuity or “grain of truth” condition of [Kalai and Lehrer \(1993\)](#), according to which rational learning must not assign zero probability to events that can occur.

The physical outcome m_t is an observable categorical variable on the support $\mathcal{M}_{\text{out}}$, where cont represents the continuation of captivity and the four causes of closure are: $j = 1$ payment of the ransom by the family, $j = 2$ death of the victim in captivity, $j = 3$ tactical rescue by the state, and $j = 4$ an exogenous residual closure channel with constant baseline intensity $\lambda_4 > 0$, which groups voluntary release by the captor without formal ransom, logistical fatigue, and administrative truncations of the sample. Conditional on the captor’s type and the material state $\mathcal{C}_t(\theta_K)$, its distribution is obtained from the effective intensities $\tilde{\lambda}_j(t \mid \theta_K)$. Denote by

$$Q_t(m \mid \mathcal{C}_t) := \mathbb{P}(m_t = m \mid \mathcal{C}_t)$$

the physical outcome transition kernel, and by G_t the measurable rule that materializes a concrete realization of m_t from that kernel and a uniform draw. The probability that captivity continues at the end of the period is

$$Q_t(\text{cont} \mid \mathcal{C}_t) = \exp\left(-\sum_{j=1}^4 \tilde{\lambda}_j(t \mid \theta_K) \Delta t\right),$$

when the episode closes, the cause is drawn with relative weight $\xi_j(t \mid \theta_K) = \tilde{\lambda}_j / \sum_{\ell=1}^4 \tilde{\lambda}_\ell$, so that causes with higher intensity concentrate the highest proba-

bility of materializing. The physical outcome m_t and the detection indicator d_t form $Y_t = (m_t, d_t) \in \mathcal{M}_{\text{out}} \times \{0, 1\}$, the projection of ζ_t onto its last two coordinates. This reduction does not restrict the state's learning, since μ_t is updated with the complete ζ_t , where (α_t^*, γ_t^*) and $\tilde{A}_t$ enter via the implementation likelihood and the material state $\mathcal{C}_t$. Its function is to isolate the stochastic draw of the competing risks block, whose central objects, intensities $\tilde{\lambda}_j$, hazards $\bar{h}_j(t)$, and causal mixture $\xi_j(t)$, are fully determined, conditional on $\mathcal{C}_t$, by the daily law of Y_t .

The modeling of risk intensities starts from distinguishing between the instantaneous intensity and its daily probabilistic projection, in line with the proportional hazards tradition of [Cox \(1972\)](#). The cause-specific intensity $\lambda_j(t)$ measures the marginal propensity to close due to cause j as a rate of occurrence per unit of time; thus, it can take any positive value and is not bounded by one. The discrete hazard $\bar{h}_j(t) \in [0, 1]$, on the other hand, is the probability that event j materializes during interval t conditional on the system having survived to the beginning of the period. The bar in $\bar{h}_j(t)$ distinguishes this marginal hazard from the public history h_t . Both objects depend on the vector of executed actions $\tilde{A}_t$ generated by the MDG process and not on the latent strategic intentions; the state's continuous instruments (α_t^*, γ_t^*) modify them through the material state $\mathcal{C}_t(\theta_K)$.

For each cause $j \in \{1, 2, 3\}$, the cause-specific intensity adopts a proportional hazards specification conditioned on the material state $\mathcal{C}_t(\theta_K)$:

$$\lambda_j(t \mid \mathcal{C}_t) = \lambda_{j0}(t) \exp(x_t' \beta_j), \quad (6)$$

where $\lambda_{j0}(t) \geq 0$ is the baseline risk of cause j , x_t is a design vector encoding the state and action space $(\theta_F, \theta_V, \theta_S, \theta_K, z, \tilde{A}_t, \alpha_t^*, \gamma_t^*, p_{\text{det},t})$, and β_j contains the latent type-specific effect $\beta_{K,j}(\theta_K)$ that captures contrasts in criminal technology

across organizations. The termination dynamics decompose into four channels: the three strategic causes $j \in \{1, 2, 3\}$ linked to institutional design and a fourth intensity $\lambda_4 > 0$ of exogenous residual closure that groups voluntary release by the captor without formal ransom, logistical fatigue, and administrative truncations of the sample; its intensity is constant baseline and invariant to type θ_K , so its likelihood factorizes orthogonally from the structural parameters in the sense of [Lancaster \(1990\)](#) and [Abbring and Berg \(2003\)](#). The maturation filter $M(t) = \min\{1, (t/T_{\text{mad}})^2\} \in [0, 1]$ scales the strategic intensities to isolate the logistical inertia of the first days of captivity. In discrete time, probabilistic closure cannot be obtained by summing individual transformations $1 - \exp(-\lambda_j)$ because that sum may exceed unity; closure is obtained by first aggregating the intensities and then projecting period survival through the aggregated exponential, in line with [Sueyoshi \(1992\)](#) and [Wooldridge \(2010\)](#). The effective intensities are:

$$\tilde{\lambda}_j(t | \mathcal{C}_t) = M(t)\lambda_j(t | \mathcal{C}_t) \quad (j = 1, 2, 3), \quad \tilde{\lambda}_4(t) = \lambda_4, \quad \Delta t = 1. \quad (7)$$

The continuation probability, the daily closure rate, and the conditional distribution of causes are obtained from:

$$p_{\text{Cont},t} = \exp\left(-\sum_{j=1}^4 \tilde{\lambda}_j(t)\right), \quad (8)$$

$$q(t) = 1 - p_{\text{Cont},t}, \quad (9)$$

$$\xi_j(t) = \frac{\tilde{\lambda}_j(t)}{\sum_{\ell=1}^4 \tilde{\lambda}_\ell(t)}, \quad j = 1, \dots, 4, \quad (10)$$

so that $\bar{h}_j(t) = q(t)\xi_j(t)$ is the marginal hazard of cause j and $\sum_j \bar{h}_j(t) = q(t) \leq 1$ by construction. The probability $\xi_j(t)$ acts as the market share of each risk and orthogonalizes two dimensions of the process: the governance of duration (the when), determined by the aggregate density of the intensities, and the governance

of the nature of the closure (the how), governed by the relative weights. It is through $\xi_j(t)$ that Bayesian learning exerts its strategic role: as μ_t concentrates on the true type, the probability mass is redistributed among exit causes, shifting the likelihood toward the outcomes that are more probable under that type, without altering the temporal scale of captivity. The probability of detecting collusion $p_{\text{det},t}$ is specified through a logistic response function conditioned on the state's instruments:

$$p_{\text{det},t}(\theta_K) = \Lambda(\eta_0(\theta_K) + \eta_1\alpha_t^* + \eta_2\gamma_t^*), \quad \Lambda(u) = \frac{1}{1 + \exp(-u)}. \quad (11)$$

where $\eta_0(\theta_K)$ is the type-specific intercept that captures the baseline detectability of the criminal organization. Because different organizations differ in their propensity to use traceable formal financial channels, d_t is informative about θ_K . The parameters $\eta_1 > 0$ and $\eta_2 > 0$ capture the marginal effect of the financial block and operational pressure. The logistic function guarantees $p_{\text{det},t}(\theta_K) \in (0, 1)$ and monotonicity increasing in (α_t^*, γ_t^*) : greater financial interdiction or operational pressure raises the probability of observing $d_t = 1$, with a baseline level that varies according to the type's concealment technology.

To operationalize the informative quality in rescue planning, the precision index ι_t is defined as the probability mass of the modal type,

$$\iota_t = \max_{\theta \in \Theta_K} \mu_t(\theta), \quad \hat{\theta}_t = \arg \max_{\theta \in \Theta_K} \mu_t(\theta), \quad (12)$$

where $\hat{\theta}_t$ is the captor's estimated profile. Let $s_t \in \{0, 1\}$ be the binary indicator of victim survival at the end of the operational margin of the period, where $s_t = 1$ means that the victim survives and $s_t = 0$ that she does not survive; this indicator operates in the rescue branch and does not redefine the categorical support of m_t .

Conditional on a focal rescue attempt, the scalar probability of survival $p_{\text{surv},t}$ is modeled through a logistic response that penalizes the misalignment between the estimated profile and the true type:

$$p_{\text{surv},t}(\iota_t, \hat{\theta}_t, \theta_K) := \Lambda(\alpha_{\text{leth}}(\theta_K) + \beta_R \cdot \iota_t \cdot \mathbf{1}\{\hat{\theta}_t = \theta_K\}), \quad (13)$$

where $\alpha_{\text{leth}}(\theta_K)$ is the intrinsic lethality of the type, $\beta_R > 0$ is the productivity of accumulated intelligence, and $\mathbf{1}\{\hat{\theta}_t = \theta_K\}$ indicates whether the state has correctly identified the captor. When the identification is correct, ι_t enhances $p_{\text{surv},t}$ according to β_R ; when there is an identification error, the accumulated intelligence becomes inoperative, and the survival probability collapses to the baseline level determined by $\alpha_{\text{leth}}(\theta_K)$.

The daily decision mechanism is projected onto a longitudinal framework via the episode survival function $S(t)$ and the cumulative incidence functions $F_j(t)$. $S(t)$ denotes the probability that the episode remains active until the end of period t , and $F_j(t)$ measures the cumulative probability of closure by cause j up to that same horizon. With $S(0) = 1$, these objects are constructed as

$$S(t) = \prod_{t'=1}^t p_{\text{Cont},t'}, \quad F_j(t) = \sum_{t'=1}^t \bar{h}_j(t') S(t' - 1), \quad S(t) + \sum_{j=1}^4 F_j(t) = 1. \quad (14)$$

The constraint $S(t) + \sum_j F_j(t) = 1$ guarantees the accounting consistency of the system: all probability not assigned to the continuation of the episode is distributed among the four causes of closure. By weighting each marginal hazard $\bar{h}_j(t')$ by the remaining survival $S(t' - 1)$, the system operates under the formal treatment of competing risks and operationalizes the state's strategic dilemma between intervening immediately or delaying action to accumulate informational capital ι_t .

The technical probability of capture is modeled via a logistic response that keeps the probability within the unit interval and allows public instruments to smoothly modify the marginal effectiveness of the intervention:

$$p_{\text{cap}}(a, \theta_i, \theta_S, \alpha_t^*, \gamma_t^*) = \Lambda(\delta_a + c_0(\theta_i) + c_\alpha(\theta_i)\alpha_t^* + c_\gamma(\theta_i)\gamma_t^* + c_S(\theta_S)), \quad (15)$$

where δ_a captures the impact of the executed intervention mode itself, $c_0(\theta_i)$ allows for baseline heterogeneity by type, $c_\alpha(\theta_i)$ and $c_\gamma(\theta_i)$ capture how the financial block and operational pressure alter the probability of capture against that specific type, and $c_S(\theta_S)$ incorporates the institutional capacity of the state. The effective probability of capture is defined as the expectation of p_{cap} over the stochastic realizations of the executed state action, conditional on the tuple $\mathcal{Q}_t^{\text{Cap}} := (a_t^{S*}, \alpha_t^*, \gamma_t^*, \theta_K, \theta_S)$:

$$\begin{aligned} \tilde{p}_{\text{cap},t}(\theta_i) &:= \mathbb{E}_{\tilde{a}_t^S | \mathcal{Q}_t^{\text{Cap}}} [p_{\text{cap}}(\tilde{a}_t^S, \theta_i, \theta_S, \alpha_t^*, \gamma_t^*)] \\ &= \sum_{a \in \mathcal{A}^S} \mathbb{P}_{\text{I},S}(a \mid a_t^{S*}, X_t) \cdot p_{\text{cap}}(a, \theta_i, \theta_S, \alpha_t^*, \gamma_t^*), \end{aligned} \quad (16)$$

where $\mathcal{Q}_t^{\text{Cap}}$ collects the optimal intention of the state, the executed instruments, and the relevant types of the episode. This formulation separates two levels: $\mathbb{P}_{\text{I},S}$ describes the probability that the state's intention is translated into a concrete action through the MDG process, and p_{cap} measures the technical probability of capture under the instruments actually executed; in this way, $\tilde{p}_{\text{cap},t}(\theta_i)$ combines implementation uncertainty and operational capacity, avoiding mechanically identifying the planned decision with the observable police outcome.

The mechanism incorporates an additional source of information through micro-level acoustic evidence and learning by absence of communication. In each period t , the state can observe a voice window or face operational silence. When there is

a signal, it is summarized in a feature vector $x_t^{obs} \in \mathbb{R}^k$ that decomposes the observed signal into the true acoustic pattern of the type ($x_t^{true}(\theta_K)$), channel noise (ε_L), and strategic camouflage (ε_S), with both perturbations being Gaussian and independent of each other:

$$x_t^{obs} = x_t^{true}(\theta_K) + \varepsilon_L + \varepsilon_S, \quad \varepsilon_L \sim \mathcal{N}(0, \Sigma_L), \quad \varepsilon_S \sim \mathcal{N}(0, \Sigma_S), \quad \varepsilon_L \perp \varepsilon_S. \quad (17)$$

The Bayesian architecture of the mechanism rests on four likelihoods that convert each observed day into evidence about the captor's type θ_K . The first captures the information contained in executed actions as realizations of the MDG process ($\mathcal{L}_{I,t}$); the second and third collect the information from the physical outcome m_t and the detection signal d_t through the competing risks block; and the fourth incorporates acoustic evidence or operational silence through the communication likelihood $\mathcal{L}_{C,t}$. Analogously to the job search models of [McCall \(1970\)](#), a continuation day updates the posterior by non-occurrence and a closure updates it through its cause, so that the model converts each observable realization into statistical information without requiring an additional observational layer. The implementation likelihood is the first of these four sources:

$$\mathcal{L}_{I,t}(\theta_K) := \prod_{i \in \{F,K,S\}} \mathbb{P}_{I,i}(\tilde{a}_t^i \mid a_t^{i*}(\theta_K), X_t), \quad (18)$$

where a_t^{K*} may depend directly on θ_K , while a_t^{F*} and a_t^{S*} depend on the public history and on μ_t ; if an executed action was unlikely under the plan induced by a type, that type receives lower posterior weight.

The likelihood of the physical outcome m_t distinguishes two regimes. On a continuation day no competing risk materializes and the absence of closure favors

types whose exit rates predict greater continuity:

$$\mathcal{L}_H^{\text{cont}}(t \mid \theta_K) = p_{\text{Cont},t|\theta_K} = \exp\left(-\sum_{j=1}^4 \tilde{\lambda}_j(t \mid \theta_K) \Delta t\right). \quad (19)$$

When the day ends due to cause $j \in \{1, 2, 3, 4\}$, the outcome identifies the type by the cause that produced it:

$$\mathcal{L}_H^{\text{out}}(t, j \mid \theta_K) = \bar{h}_j(t \mid \theta_K). \quad (20)$$

Let $g : \{1, 2, 3, 4\} \rightarrow \mathcal{M}_{\text{term}}$ be the bijection $g(1) = \text{pay}$, $g(2) = \text{kill}$, $g(3) = \text{res}$, $g(4) = \text{rel}$. Both regimes are written as a single expression:

$$\mathbb{P}(m_t \mid \theta_K, \mathcal{C}_t) = p_{\text{Cont},t|\theta_K}^{\mathbf{1}\{m_t=\text{cont}\}} \prod_{j=1}^4 \bar{h}_j(t \mid \theta_K)^{\mathbf{1}\{m_t=g(j)\}}, \quad (21)$$

a five-cell multinomial (cont and $\mathcal{M}_{\text{term}}$) where g converts index j into the label of m_t .

Regardless of the physical outcome, the detection signal d_t provides a second source of information with Bernoulli likelihood:

$$\mathbb{P}(d_t \mid \theta_K, \alpha_t^*, \gamma_t^*) = p_{\text{det},t}(\theta_K)^{d_t} (1 - p_{\text{det},t}(\theta_K))^{1-d_t}. \quad (22)$$

The mechanism incorporates a third source of information through micro-level acoustic evidence and learning by absence of communication. Each type θ_K possesses a reference profile $(\bar{x}(\theta_K), \Sigma(\theta_K))$; the diagonal Gaussian likelihood assigns greater weight to types whose profiles make the observed signal more plausible:

$$\mathcal{L}_{\text{voz},t}(\theta_K) \propto \exp\left(-\frac{1}{2} \sum_{i=1}^k \frac{(x_{t,i}^{\text{obs}} - \bar{x}_i(\theta_K))^2}{\tilde{\sigma}_i(\theta_K)^2}\right). \quad (23)$$

Let $V_t \in \{0, 1\}$ be the indicator of voice emission at t and $\pi_{\text{call}}(\theta_K) \in (0, 1)$ the contact probability of each type, interpreted as logistical urgency; both the call and silence are informative. The type-conditional communication likelihood is

$$\mathcal{L}_{C,t}(\theta_K | V_t) = \begin{cases} [\mathcal{L}_{\text{voz},t}(\theta_K) \pi_{\text{call}}(\theta_K)]^{\omega_{\text{voz}}}, & V_t = 1, \\ [1 - \pi_{\text{call}}(\theta_K)]^{\omega_{\text{voz}}}, & V_t = 0, \end{cases} \quad (24)$$

where $\omega_{\text{voz}} \in [0, 1]$ is the learning weight; when there is no active information, $\omega_{\text{voz}} = 0$ and $\mathcal{L}_{C,t} \equiv 1$.

The three sources, executed actions, physical outcome, and detection signal, come from distinct operational margins connected by the same material state, which allows factoring the joint likelihood of the physical-institutional block:

$$\mathcal{L}_{F,t}(\theta_K) = \underbrace{\mathcal{L}_{I,t}(\theta_K)}_{\text{executed actions}} \cdot \underbrace{\mathbb{P}(m_t | \theta_K, \mathcal{C}_t)}_{\text{competing risks}} \cdot \underbrace{\mathbb{P}(d_t | \theta_K, \alpha_t^*, \gamma_t^*)}_{\text{detection}}. \quad (25)$$

If continuity and positive detection are observed, the risks block reduces the weight of types with high hazards and the detection block reduces that of types with low probability of fiscalization; if rescue and absence of detection are observed, the competitive core favors types with higher rescue hazard; $\mathcal{L}_{I,t}$ penalizes types for which the executed triplet was improbable under the MDG. The unified Bayes rule combines the physical-institutional and communication blocks:

$$\mu_{t+1}(\theta_K) = \frac{\mu_t(\theta_K) \mathcal{L}_{F,t}(\theta_K) \mathcal{L}_{C,t}(\theta_K | V_t)}{\sum_{\tilde{\theta} \in \Theta_K} \mu_t(\tilde{\theta}) \mathcal{L}_{F,t}(\tilde{\theta}) \mathcal{L}_{C,t}(\tilde{\theta} | V_t)}. \quad (26)$$

The degenerate rule illustrates the modularity of the system: when $\omega_{\text{voz}} = 0$ or there is no voice data, $\mathcal{L}_{C,t} \equiv 1$ and updating relies only on the physical-institutional record $(\tilde{A}_t, m_t, d_t)$; if implementation evidence is additionally set aside by fixing $\mathcal{L}_{I,t} \equiv 1$, updating reduces to the minimal record (m_t, d_t) . The

principle governing which objects enter the posterior is simple: the Bayes operator can update only on realizations publicly observable at the close of the period. Therefore, strategic intentions a_t^{i*} do not enter directly; only their observable realizations $\tilde{A}_t$ do so through $\mathcal{L}_{I,t}$; and capture probabilities $\tilde{p}_{\text{cap},t}$ and survival probabilities $p_{\text{surv},t}$ are already contained in the law of the physical outcome $\mathbb{P}(m_t \mid \theta_K, \mathcal{C}_t)$, so incorporating them separately would imply double counting. Information about θ_K is not introduced as an exogenous signal but as an observable product of a strategic optimization process subjected to the MDG filter, so that, under learning with increasing information in the sense of [Blackwell and Dubins \(1962\)](#), the sequence of realizations progressively filters incompatible types.

2.4. Effective probabilities

In all decision branches, the probabilities that weight the outcomes are defined as conditional expectations over the realizations of the MDG process induced by the conditioning tuple $\mathcal{Q}_t$ pertinent to each action. This separation between the underlying technological probability and its stochastic realization avoids mechanically identifying the strategic intention with the observable result. The four effective probabilities that feed the Bellman objects are:

$$\tilde{p}_{\text{cap},t}(\theta_i) := \mathbb{E}_{\tilde{a}_t^S \mid \mathcal{Q}_t^{\text{Cap}}} [p_{\text{cap}}(\tilde{a}_t^S, \theta_i, \theta_S, \alpha_t^*, \gamma_t^*)], \quad (27)$$

$$\tilde{p}_{\text{pay},t}(\theta_i) := \mathbb{E}_{\tilde{A}_t \mid \mathcal{Q}_t^{\text{Cont}}} [\mathbb{P}(m_t = \text{pay} \mid \tilde{A}_t, X'_t, \theta_K = \theta_i)], \quad (28)$$

$$\tilde{p}_{\text{surv},t}(\theta_i) := \mathbb{E}_{\tilde{A}_t \mid \mathcal{Q}_t^{\text{Coop}}} [p_{\text{surv},t}(\iota_t, \hat{\theta}_t, \theta_K = \theta_i)], \quad (29)$$

$$\tilde{p}_{\text{rel},t}(\theta_i) := \mathbb{E}_{\tilde{A}_t \mid \mathcal{Q}_t^{\text{Col}}} [\mathbb{P}(m_t = \text{rel} \mid \tilde{A}_t, R, \theta_K = \theta_i)], \quad (30)$$

The conditioning tuples specify the branch action and the executed instruments for each expectation. For capture, $\mathcal{Q}_t^{\text{Cap}} := (a_t^{S*}, \alpha_t^*, \gamma_t^*, \theta_K, \theta_S)$. The three

remaining branches share the structure $\mathcal{Q}_t^b := (a_t^{S*}, a_t^{b*}, a_b, \alpha_t^*, \gamma_t^*, \mu_t, \theta_K)$, where $b \in \{\text{Cont}, \text{Coop}, \text{Col}\}$ and a_b is the corresponding branch action (a_{cont} , a_{coop} , or a_{col}). In this way, the model maintains coherence between the chosen strategic branch and the evaluated stochastic realization.

2.5. Rational behavior and best responses

Captor's behavior.

In each period t , the captor solves an optimal stopping problem, selecting the action $a_t^K \in \mathcal{A}^K = \{a_{\text{rel}}, a_{\text{kill}}, a_{\text{cont}}\}$ that maximizes its value function given its private type θ_K , the public history h_t , and the institutional environment θ_S . Comparing the utilities of releasing, killing, and continuing captivity, its best response in compact form is

$$V_t^K(\theta_K) = \max\{U_{\text{rel},t}^K(\theta_K), U_{\text{kill},t}^K(\theta_K), V_{\text{cont},t}^K(\theta_K)\}. \quad (31)$$

In the death branch, the payoff weights the reputational benefit $\eta(\theta_K)$ against the expected penalty for capture:

$$U_{\text{kill}}^K(\theta_K, \theta_S) = (1 - \tilde{p}_{\text{cap},t}(\theta_K))\eta(\theta_K) - \tilde{p}_{\text{cap},t}(\theta_K) F_{\text{cap}}(\theta_K, \theta_S). \quad (32)$$

In the release branch, the captor incurs the disutility of yielding without obtaining rent, reflected in lost reputation or foregone income:

$$U_{\text{rel}}^K(\theta_K) = -\kappa_{\text{rel}}(\theta_K), \quad \kappa_{\text{rel}}(\theta_K) \geq 0. \quad (33)$$

Continuation combines expected ransom net of financial blockade, operating cost, expected capture penalty, and discounted future value by type. Since θ_K is private information of the captor, μ_t is not an uncertainty for it but a strategic object:

the evolution of μ_{t+1} conditions the state's future policy and forms part of its intertemporal calculation.

$$\begin{aligned}
V_{\text{cont},t}^K(\theta_K) = & \tilde{p}_{\text{pay},t}(\theta_K) R(1 - \alpha_t^*) - C_t(\gamma_t^*, \theta_K) \\
& - \tilde{p}_{\text{cap},t}(\theta_K) F_{\text{cap}}(\theta_K, \theta_S) \\
& + \tilde{\beta}(\theta_K)(1 - \tilde{p}_{\text{cap},t}(\theta_K)) \mathbb{E}[V_{t+1}^K(\mu_{t+1}(a_{\text{cont}}))],
\end{aligned} \tag{34}$$

where $C_t(\gamma_t^*, \theta_K)$ is a convex cost function that is increasing in γ_t^* and varies asymmetrically across types, generating the single-crossing property that renders the persistence of captivity an informative signal for the state's Bayesian updating; $\tilde{\beta}(\theta_K) \in (0, 1)$ is the type-specific discount factor; $\tilde{p}_{\text{pay},t}$ and $\tilde{p}_{\text{cap},t}$ are the effective probabilities from Section 2.4; and the continuation of future value is taken over the evolution of the belief μ_{t+1} induced by the action a_{cont} . The optimal action of the captor a_t^{K*} arises from comparing the three branches:

$$a_t^{K*} = \arg \max_{a_t^K \in \mathcal{A}^K} \mathcal{U}_t^K(a_t^K \mid h_t, \theta_K, \theta_S), \quad \mathcal{U}_t^K = \begin{cases} U_{\text{rel}}^K(\theta_K), & \text{if } a_t^K = a_{\text{rel}}, \\ U_{\text{kill}}^K(\theta_K, \theta_S), & \text{if } a_t^K = a_{\text{kill}}, \\ V_{\text{cont},t}^K(\theta_K), & \text{if } a_t^K = a_{\text{cont}}. \end{cases} \tag{35}$$

Family's behavior:

In each period t , the family chooses $a_t^F \in \mathcal{A}^F = \{a_{\text{coop}}, a_{\text{col}}\}$ by maximizing its expected utility over the uncertainty regarding the captor's type, weighted by its private information set $\mathcal{I}_t^F$. In the cooperation branch, the family follows the institutional path, paying the cooperation effort cost $e_t(\gamma_t, \theta_F)$, which captures the frictions and risk of retaliation under interdiction, and is convex and increasing in γ_t with asymmetric variation across family types that generates a single-crossing property analogous to that of the captor; in return, it obtains the expected value of

the victim's survival weighted by $V_L > 0$:

$$\mathcal{U}_t^F(a_{\text{coop}} \mid \theta_F) = \mathbb{E}_{\theta_K \mid \mathcal{I}_t^F}[\tilde{p}_{\text{surv},t}(\theta_K)] V_L - e_t(\gamma_t, \theta_F). \quad (36)$$

In the collusion branch, the family clandestinely pays the ransom R to secure the victim's release, but discounts the probability of detection by the state and the fine F_{col} :

$$\mathcal{U}_t^F(a_{\text{col}} \mid \theta_F) = \mathbb{E}_{\theta_K \mid \mathcal{I}_t^F}[\tilde{p}_{\text{rel},t}(\theta_K)] V_L - R - p_{\text{det},t} F_{\text{col}}, \quad (37)$$

where $\tilde{p}_{\text{surv},t}$ and $\tilde{p}_{\text{rel},t}$ are the effective probabilities from Section 2.4. The optimal action of the family arises from comparing the two branches:

$$a_t^{F*} = \arg \max_{a_t^F \in \mathcal{A}^F} \mathcal{U}_t^F(a_t^F \mid \theta_F), \quad \mathcal{U}_t^F = \begin{cases} \mathcal{U}_t^F(a_{\text{coop}}), & \text{if } a_t^F = a_{\text{coop}}, \\ \mathcal{U}_t^F(a_{\text{col}}), & \text{if } a_t^F = a_{\text{col}}. \end{cases} \quad (38)$$

State's behavior.

In each period t , given the public history $\mathcal{F}_t^{\text{pub}}$ and the belief vector $\mu_t \in \Delta(\Theta_K)$, the state solves a Bayesian dual control problem by choosing the discrete action $a_t^S \in \{a_{\text{res}}, a_{\text{neg}}\}$ and the continuous instruments $(\alpha_t, \gamma_t) \in [0, 1]^2$. Its objective is not to maximize its own utility but to minimize an expected social loss that combines human risk, economic costs, operational costs, and an informational motive. The conditional loss by type varies by branch:

$$L_t(a_t^S, \alpha_t, \gamma_t, \theta_K, \iota_t) = \begin{cases} V_t^R(\iota_t, \hat{\theta}_t, \theta_K, \alpha_t, \gamma_t), & \text{if } a_t^S = a_{\text{res}}, \\ V_t^N(\theta_K, \alpha_t, \gamma_t), & \text{if } a_t^S = a_{\text{neg}}, \end{cases} \quad (39)$$

where under tactical rescue the state incurs the cost of a lethal failure of the operation and the cost of deployment, $V_t^R := \omega_k[1 - p_{\text{surv},t}(\iota_t, \hat{\theta}_t, \theta_K)] + C_{\text{ops}}(\gamma_t, \alpha_t; \theta_K)$, and under negotiation it incurs the cost of net criminal transfer, the lethality of captivity, and the maintenance of the containment ring, $V_t^N := \omega_p R(1 - \alpha_t) + \omega_k \bar{h}_2(t \mid$

$\theta_K, \mathcal{C}_t) + C_{\text{maint}}(\gamma_t, \alpha_t; \theta_K)$; here $\bar{h}_2(t \mid \theta_K, \mathcal{C}_t) = q(t \mid \theta_K) \xi_2(t \mid \theta_K)$ is the marginal hazard of death from captivity, derived from the competing risks system (10), and the weights $\omega_k > 0$ and $\omega_p > 0$ calibrate the human loss and economic transfer respectively. The cost functions C_{ops} and C_{maint} are quadratic, continuous, and convex in (α_t, γ_t) , with parameters that depend on θ_K to capture logistical, territorial, and financial heterogeneity; the cross terms in $\gamma_t \alpha_t$ capture complementarity or substitutability between financial block and operational pressure depending on the type of criminal organization faced. The state's executable problem is

$$(a_t^{S*}, \alpha_t^*, \gamma_t^*) \in \arg \min_{(a_t^S, \alpha_t, \gamma_t) \in \Gamma_t(\mu_t)} \left\{ \sum_{\theta \in \Theta_K} \mu_t(\theta) L_t(a_t^S, \alpha_t, \gamma_t, \theta, \iota_t) \right. \\ \left. + \Pi_{t,a^S}^S(\alpha_t, \gamma_t; \mu_t) - \psi_H \Delta H_t(\alpha_t, \gamma_t) \right\}, \quad (40)$$

where $\Gamma_t(\mu_t)$ incorporates the belief-weighted constraints IC^K , IR^K , and IR^F . To prevent the state from choosing arbitrarily extreme instruments under uncertainty, type- and branch-specific counterfactual benchmarks are constructed to indicate what the state would do if it knew θ with certainty: $(\alpha_{t,b}^{\text{bench}}(\theta), \gamma_{t,b}^{\text{bench}}(\theta)) \in \arg \min_{(\alpha, \gamma)} V_t^b(\theta, \alpha, \gamma)$ for $b \in \{R, N\}$. The Bayesian center of branch b weights these benchmarks by the current belief, $(\alpha_{t,b}^\mu, \gamma_{t,b}^\mu) := \sum_{\theta} \mu_t(\theta) (\alpha_{t,b}^{\text{bench}}(\theta), \gamma_{t,b}^{\text{bench}}(\theta))$, and the penalty $\Pi_{t,b}^S$ measures how much the candidate policy deviates from this center:

$$\Pi_{t,b}^S(\alpha_t, \gamma_t; \mu_t) = \chi_\alpha (\alpha_t - \alpha_{t,b}^\mu)^2 + \chi_\gamma (\gamma_t - \gamma_{t,b}^\mu)^2, \quad \chi_\alpha, \chi_\gamma \geq 0, \quad (41)$$

which allows exploring alternative instruments but charges a premium when the instrument moves away from the Bayesian center. The informational motive $-\psi_H \Delta H_t$ operationalizes active exploration: the state recognizes that (α_t, γ_t)

modifies risk intensities and the probability of detection, so a policy may be valuable not only for reducing current material loss but for separating likelihoods and reducing future uncertainty, in line with the dynamic acquisition of information of [Zhong \(2022\)](#). The expected information gain is the expected reduction in the entropy of [Shannon \(1948\)](#) over the joint public signal (m_t, d_t) ,

$$\Delta H_t(\alpha_t, \gamma_t) = H(\mu_t) - \sum_m \sum_{d \in \{0,1\}} \mathbb{P}(m_t = m, d_t = d \mid \mu_t, \alpha_t, \gamma_t, \mathcal{C}_t) H(\mu_{t+1}^{m,d}), \quad (42)$$

where $\mu_{t+1}^{m,d}$ is the counterfactual posterior that would result from observing the pair (m, d) under the candidate policy, and the joint probability factorizes by conditional independence given the type: $\mathbb{P}(m, d \mid \mu_t, \alpha_t, \gamma_t, \mathcal{C}_t) = \sum_{\theta} \mu_t(\theta) \mathbb{P}(m_t = m \mid \theta, \mathcal{C}_t) \mathbb{P}(d_t = d \mid \theta, \alpha_t, \gamma_t)$. The voice channel V_t and the implementation signal $\tilde{A}_t$ are excluded from ΔH_t : the former does not depend on (α_t, γ_t) and is modeled separately in $\mathcal{L}_{C,t}$; the latter captures implementation friction, not type discrimination. If $\psi_H = 0$, the state is purely myopic; if $\psi_H > 0$, policies that buy learning receive a discount on the loss; when μ_t concentrates on the true type, ΔH_t falls and endogenously shuts down the exploratory motive. The solution is obtained by partition: first, the continuous instruments are minimized within each branch, incorporating the penalty and the entropic motive, and then the optimal

floors are compared. The loss floors of each branch are

$$\begin{aligned} \tilde{V}_t^{R,*}(\iota_t, \hat{\theta}_t, \mu_t) = & \min_{(\alpha_t, \gamma_t) \in \Gamma_t^R(\mu_t)} \left\{ \sum_{\theta \in \Theta_K} \mu_t(\theta) V_t^R(\iota_t, \hat{\theta}_t, \theta, \alpha_t, \gamma_t) \right. \\ & \left. + \Pi_{t,R}^S(\alpha_t, \gamma_t; \mu_t) - \psi_H \Delta H_t(\alpha_t, \gamma_t) \right\}, \end{aligned} \quad (43)$$

$$\begin{aligned} \tilde{V}_t^{N,*}(\mu_t) = & \min_{(\alpha_t, \gamma_t) \in \Gamma_t^N(\mu_t)} \left\{ \sum_{\theta \in \Theta_K} \mu_t(\theta) V_t^N(\theta, \alpha_t, \gamma_t) \right. \\ & \left. + \Pi_{t,N}^S(\alpha_t, \gamma_t; \mu_t) - \psi_H \Delta H_t(\alpha_t, \gamma_t) \right\}. \end{aligned} \quad (44)$$

so that each floor integrates the expected material loss, the penalty for deviation from the Bayesian center, and the informational subsidy. The state's optimal social cost is

$$\mathcal{L}_t^{S*} = \min\{\tilde{V}_t^{R,*}(\iota_t, \hat{\theta}_t, \mu_t), \tilde{V}_t^{N,*}(\mu_t)\}. \quad (45)$$

If $(\alpha_{\text{res}}^*, \gamma_{\text{res}}^*)$ and $(\alpha_{\text{neg}}^*, \gamma_{\text{neg}}^*)$ denote the continuous argminima of the rescue and negotiation branches, the optimal executed policy is

$$(a_t^{S*}, \alpha_t^*, \gamma_t^*) = \begin{cases} (a_{\text{res}}, \alpha_{\text{res}}^*, \gamma_{\text{res}}^*), & \text{if } \tilde{V}_t^{R,*}(\iota_t, \hat{\theta}_t, \mu_t) \leq \tilde{V}_t^{N,*}(\mu_t), \\ (a_{\text{neg}}, \alpha_{\text{neg}}^*, \gamma_{\text{neg}}^*), & \text{if } \tilde{V}_t^{R,*}(\iota_t, \hat{\theta}_t, \mu_t) > \tilde{V}_t^{N,*}(\mu_t). \end{cases} \quad (46)$$

In case of numerical equivalence between the loss floors, rescue prevails due to the weak inequality of the first condition.

Individual rationality and incentive compatibility (IR/IC).

The minimization of the state's social loss is bounded by three design constraints that guarantee the viability of learning, the deterrence of crime, and civil cooperation. Since the state operates under incomplete information, IC^K and IR^K are formulated as expected-value constraints under the posterior belief μ_t , in line with Bayesian mechanism design of [Holmström and Myerson \(1983\)](#). This formulation does not require mechanical type-by-type separation in every period,

but imposes the constraint on the prevailing type distribution; it converges to the standard condition for the true type when $\mu_t \rightarrow \delta_{\theta_K^*}$.

The first is the *Bayesian incentive compatibility of the kidnapper* (IC^K), which requires that, in expected value over the true type, the trajectory assigned to each type dominates the mimicry of any rival type:

$$\mathbb{E}_{\theta \sim \mu_t} [V_t^K(a^*(\theta) \mid \theta, \alpha_t, \gamma_t) - V_t^K(a^*(\theta_j) \mid \theta, \alpha_t, \gamma_t)] \geq 0, \quad \forall \theta_j \in \Theta_K. \quad (47)$$

When $\mu_t(\theta^*) \rightarrow 1$, the constraint collapses into the incentive compatibility of the true type against each rival mimicry.

The second is the *Bayesian participation constraint of the kidnapper* (IR^K), which operates under a logic of strategic deterrence: the reservation utility linked to peaceful release, $U_{\text{rel}}^K(\theta) = -\kappa_{\text{rel}}(\theta)$, must dominate in expected value the best criminal alternative between continuation and executing the hostage:

$$\mathbb{E}_{\theta \sim \mu_t} [U_{\text{rel}}^K(\theta) - \max\{V_{\text{cont},t}^K(\theta, \alpha_t, \gamma_t), U_{\text{kill}}^K(\theta, \theta_S)\}] \geq 0. \quad (48)$$

The third is the *participation constraint of the family* (IR^F), which disciplines the financing channel by requiring that the expected utility of cooperating with the state dominates the return of colluding via clandestine payment of the ransom:

$$\mathcal{U}_t^F(a_{\text{coop}}) \geq \mathcal{U}_t^F(a_{\text{col}}). \quad (49)$$

The three constraints jointly define the feasible set $\Gamma_t(\mu_t)$ within which the state optimizes, and anchor the functional form of the likelihood by restricting the space of admissible strategic profiles in PBE.

2.6. Type Identification Theorem for Rational Types in Unknown Environments

Bayesian convergence does not imply identification. As [Doob \(1953\)](#) and [Blackwell and Dubins \(1962\)](#) show, beliefs can stabilize without concentrating

on the true type when different profiles generate observationally equivalent signals. In the sense of [Kalai and Lehrer \(1993\)](#), Bayesian learning may suffice to predict the future path correctly and stop at a subjective equilibrium; analogously, the self-confirming equilibrium of [Fudenberg and Levine \(1993\)](#) allows beliefs to be correct on the observed trajectory and wrong on unvisited contingencies. Therefore, predictive coherence on the path does not necessarily identify the true type. This section establishes when the public history of the mechanism discriminates types in the limit, combining Bayesian coherence, statistical separation between signal laws, and the singularity criterion of [Kakutani \(1948\)](#). The strategy requires three ingredients: a public filtration over which μ_t is a martingale; a long-run identification condition that prevents distinct types from explaining the same observable trajectory; and uniform separation in total variation provided by heterogeneity in $\beta_{K,j}(\theta_K)$, which induces mutual singularity of trajectory laws and collapse of posterior mass on false types.

The public signal ζ_t is the central object that feeds the filtration. Once the PBE is fixed, $\mathbb{P}_E$ is the ex-ante law on Ω that describes the entire experiment before knowing the captor's type; the type is realized according to the structural prior $\theta_K \sim \mu_0$, and conditional on this type, the public policy, executed actions, outcomes, and public history h_t unfold. For each counterfactual type $\theta \in \Theta_K$, $\mathbb{P}_{C,\theta}$ denotes the corresponding type-conditional law on Ω under the same fixed PBE.

Definition 1 (Public signal and filtration). (i) Public signal ζ_t . The public signal ζ_t is the random vector that summarizes what is observable at the end of period t . The order of the components reproduces the way in which the public history incorporates, at the end of each period, the chosen instruments, the executed ac-

tions, and the observed results: $\alpha_t \in [0, 1]$ and $\gamma_t \in [0, 1]$ are the continuous instruments of the state; $(\tilde{a}_t^F, \tilde{a}_t^K, \tilde{a}_t^S)$ are the executed discrete actions that appear as public observables after the MDG process and do not replace the equilibrium intentions a_t^{i*} in the record; $m_t \in \mathcal{M}_{\text{out}}$ is the physical outcome ($\mathcal{M}_{\text{out}}$ is the physical outcome space defined in Section 2); and $d_t \in \{0, 1\}$ is the detection signal.

(ii) Value space. ζ_t takes values in the product space $\mathcal{Z}_t \subseteq \mathbb{R}^2 \times \mathcal{A}^F \times \mathcal{A}^K \times \mathcal{A}^{S, \text{disc}} \times \mathcal{M}_{\text{out}} \times \{0, 1\}$, where the discrete components are symbols in finite sets and the continuous coordinates are real. Thus, ζ_t encodes what is observable at the close of t and summarizes the public innovation that feeds the next history of the episode.

(iii) Public filtration $\{\mathcal{F}_t^{\text{pub}}\}_{t \geq 0}$. The initial instant is fixed by $\mathcal{F}_0^{\text{pub}} := \sigma(z, \theta_V)$, i.e., the σ -algebra generated by the observable geographic location at the opening of the episode. For each integer $t \geq 1$, $\mathcal{F}_t^{\text{pub}} := \sigma(z, \zeta_0, \dots, \zeta_{t-1})$, i.e., the σ -algebra generated by the initial location and by the signals observed strictly before starting period t (equivalent to “at the opening of day t ” in the temporal reading of the document).

(iv) Measurability and measurable structure in $\mathcal{Z}_t$. It is assumed that each $\zeta_t : (\Omega, \mathcal{F}) \rightarrow (\mathcal{Z}_t, \mathcal{B}(\mathcal{Z}_t))$ is measurable and that z is observable at the opening of the episode, so that $\sigma(z, \zeta_0, \dots, \zeta_{t-1})$ is well defined as a sub- σ -algebra of $\mathcal{F}$.

□

Once the public filtration is fixed, each counterfactual type induces its own law on the daily signal. The following definition formalizes this law as a predictive kernel, the object that the Bayes operator reweights in each period and that connects the space of latent types with the space of observables.

Definition 2 (Predictive likelihood under a counterfactual type). Fixing a PBE with profile (χ, s^F, s^K) , the stochastic dynamics of the mechanism, and a counterfactual type $\theta \in \Theta_K$, the predictive likelihood of period t is defined by

$$\ell_t(\theta; J) := \mathbb{P}_{C,\theta}(\zeta_t \in J \mid \mathcal{F}_t^{\text{pub}}), \quad J \in \mathcal{B}(\mathcal{Z}_t), \quad (50)$$

where $\mathcal{Z}_t$ is the space of possible values of the public signal and J represents an observable subset of that space. The conditioning is interpreted as the conditional expectation of $\mathbf{1}\{\zeta_t \in J\}$ given the public information accumulated up to t , defined up to $\mathbb{P}_{C,\theta}$ -null sets through standard regular versions. The likelihoods $\mathcal{L}_{F,t}$ and $\mathcal{L}_{C,t}$ are the operational realization of this kernel for the physical block and the communication block of the period.

□

The model architecture postulates that the likelihood of observable trajectories, denoted by $\mathcal{L}_t(\theta) = \mathcal{L}_{F,t}(\theta)\mathcal{L}_{C,t}(\theta)$, is not a statistically invariant object, but is subject to the Lucas Critique of [Lucas \(1976\)](#); under any modification in the mechanism χ , whether due to the tightening of instruments or a more aggressive operational stance of the state (a_t^S), agents reoptimize their strategies in response to the new institutional environment, thereby recomposing the equilibrium map (χ, s^F, s^K) and the predictive likelihood of the behaviors encoded in the informative stream ζ_t . In this framework, the conditional kernel ℓ_t acts as the formal bridge between the space of latent states ($\theta \in \Theta_K$) and the space of observables ($\mathcal{Z}_t$), where for each hypothetical type θ , the measure $\ell_t(\theta; J)$ quantifies the probability distribution of the daily signal over the Borel sets $J \in \mathcal{B}(\mathcal{Z}_t)$, defining a differentiated predictive model that allows the Bayes operator to reweight the posterior belief $\mu_{t+1}(\theta)$ after each realization of ζ_t . It is crucial to note that the

participation (IR) and incentive compatibility (IC) constraints, in the sense of [Myerson \(1979\)](#) and [Bolton and Dewatripont \(2005\)](#), do not replace this process, but rather anchor the functional form of ℓ_t by restricting the space of admissible strategic profiles; only under a profile that satisfies these conditions does the observed likelihood coincide with that induced by an implementable mechanism, thereby guaranteeing that the informational update is numerically stable, structurally identified, and fully consistent with PBE.

The belief process is represented as a martingale under the ex-ante law. The following definition formalizes the consistency of this process with the PBE, a necessary condition for the Bayesian update to depend not on the specified model but on the structural experiment of the mechanism.

Definition 3 (Beliefs and PBE-consistent Bayesian updating). A sequence of public beliefs about the captor's type must coincide with the Bayesian updating induced by the same PBE that generates the observed signals. Let $\mu_0 \in \Delta(\Theta_K)$ be the common prior over θ_K , with full conditional support on the initial public information: $\mu_0(\theta \mid z, \theta_V) > 0$ for all $\theta \in \Theta_K$. This prior summarizes the initial belief given the public location z and the victim's profile θ_V , so that

$$\mu_0(\theta \mid z, \theta_V) = \mathbb{P}_E(\theta_K = \theta \mid z, \theta_V),$$

in accordance with the formalization of the structural prior in [Section 2](#). Given a PBE, the strategic profile simultaneously determines the state's public policy, private strategies, action materialization through the MDG process, outcomes, and the ex-ante law $\mathbb{P}_E$ over public trajectories. For $t \geq 1$, let $\mu_t \in \Delta(\Theta_K)$ denote the posterior obtained by iterated Bayes from μ_0 , using the location z , the signals already observed $(\zeta_0, \dots, \zeta_{t-1})$, and the predictive kernels $\ell_u(\theta'; J)$ induced

by that same PBE. Therefore, updating with those realizations is equivalent to conditioning on the public filtration available at the opening of period t :

$$\mathcal{F}_t^{\text{pub}} = \sigma(z, \zeta_0, \dots, \zeta_{t-1}),$$

The process $\{\mu_t\}_{t \geq 0}$ is said to be consistent with the PBE if, under the ex-ante law $\mathbb{P}_E$ induced by the same equilibrium,

$$\mu_t(\theta) = \mathbb{P}_E(\theta_K = \theta \mid \mathcal{F}_t^{\text{pub}}) \quad \forall \theta \in \Theta_K, \forall t \geq 1, \quad (51)$$

where the equality is interpreted $\mathbb{P}_E$ -almost surely upon fixing regular versions of the conditional probabilities. It is assumed that $\theta_K : (\Omega, \mathcal{F}) \rightarrow (\Theta_K, \mathcal{B}(\Theta_K))$ is measurable and that $\{\theta\} \in \mathcal{B}(\Theta_K)$ for all $\theta \in \Theta_K$; in particular, if Θ_K is finite, it suffices to take the power set σ -algebra. Thus, $\mathbf{1}\{\theta_K = \theta\}$ is measurable and integrable under $\mathbb{P}_E$, and the posterior μ_t coincides with the conditional marginal of the type given the observed public history. Furthermore, μ_0 coincides with the initial marginal of θ_K once conditioned on (z, θ_V) , before the signals of the episode refine the structural prior.

□

Definition 3 implies that $\mu_t(\theta)$ is a bounded martingale under $\mathbb{P}_E$: the current belief is the unbiased prediction of the future belief given the accumulated public history. The following lemma formalizes the almost sure convergence that follows, relying on the convergence theorem of Doob (1953). Its proof is presented in Appendix C.

Lemma 4 (Asymptotic stability of beliefs). *Under the ex-ante law $\mathbb{P}_E$, the belief process $\{\mu_t(\theta)\}_{t \geq 0}$ is a bounded martingale that converges $\mathbb{P}_E$ -almost surely to a limit random variable Z_θ with $Z_\theta \in [0, 1]$ $\mathbb{P}_E$ -a.s.*

□

To establish the asymptotic properties of the learning process under the true probability measure $\mathbb{P}_{\theta_0}$, a fundamental tension must be resolved regarding the conditional independence of the observable history. In the live perfect Bayesian equilibrium, the planner’s active dual control couples the policy instruments sequentially with past realized signals, as in [Aghion et al. \(1991\)](#) and [Easley and Kiefer \(1988\)](#). However, structural identification can be formally established by evaluating the system’s filtering capacity over a frozen, deterministic policy path, following [Blackwell and Dubins \(1962\)](#). This conceptual device is formalized here under the notion of an informational slicing of the game tree, deferring the formal measure-theoretic proofs and regularities to Appendix C, where the mutual singularity of infinite product measures is derived via the product-measure criterion of [Kakutani \(1948\)](#).

Before defining the asymptotic event, it is useful to separate two planes of the model. Following the distinction between structure and measurement in dynamic estimation of [Rust \(1987\)](#) and [Magnac and Thesmar \(2002\)](#), the measures $\mathbb{P}_{C,\theta}$ used in the following definitions and in the Identification Theorem refer to the strategic competing-risks block, $j \in \{1, 2, 3\}$. In that domain, $\lambda_4 = 0$ is fixed, so the residual closure channel does not determine asymptotic type identification. In the data and in calibration, by contrast, $\lambda_4 > 0$ captures type-invariant residual closures; its contribution is therefore treated as an orthogonal component of the likelihood, in the sense of [Lancaster \(1990\)](#) and [Abbring and Berg \(2003\)](#). This separation allows the singularity criterion of [Kakutani \(1948\)](#) to operate on the signals generated by strategic risks. For the convergence of Lemma 4 to concentrate on the true type rather than a diffuse limit, the analysis is restricted to the event in which captivity generates an infinite sequence of informative public

signals.

Definition 5 (Prolonged captivity event). Let τ be the stopping time defined in (3). The prolonged captivity event is defined as

$$\mathcal{T} := \{\tau = \infty\} = \{\omega \in \Omega : m_t(\omega) \in \mathcal{M}_{\text{cont}} \forall t \geq 0\}.$$

In $\mathcal{T}$, the episode does not admit a terminal outcome in any finite period and generates an infinite sequence of public signals $\{\zeta_t\}_{t \geq 0}$.

□

Asymptotic identification requires that types other than the true type cannot sustain the same signal trajectory in $\mathcal{T}$, either because they cannot maintain captivity indefinitely or because their signal laws are mutually singular with those of the true type.

Definition 6 (Asymptotic identification in the event $\mathcal{T}$). The model is asymptotically identifiable in $\mathcal{T}$ if, for all $\theta \in \Theta_K$ with $\theta \neq \theta_K^*$, at least one of the following conditions is satisfied:

- (a) Exclusion by survival. $\mathbb{P}_{C,\theta}(\mathcal{T}) = 0$.
- (b) Tail singularity of signals. $\mathbb{P}_{C,\theta}(\mathcal{T}) > 0$ and the laws of $\{\zeta_t\}_{t \geq 0}$ restricted to $\mathcal{T}$ under $\mathbb{P}_{C,\theta}$ and $\mathbb{P}_{C,\theta_K^*}$ are mutually singular.

□

The heterogeneity in the risk coefficients $\beta_{K,j}(\theta_K)$ is the structural channel that operationalizes the separation of Definition 6: different criminal organizations generate distinct survival and detection signatures in total variation, preventing sustained observational equivalence and activating the singularity criterion of

[Kakutani \(1948\)](#). The state’s ability to distinguish types in the asymptotic limit does not depend on a specific functional form, but on a fundamental property of information theory: any family of predictive models that guarantees recurrent separation in total variation between its signal laws will induce mutual singularity of the trajectory measures and the consequent concentration of Bayesian beliefs. The following lemma specializes this principle to the case of competing risks, demonstrating how the architecture of [Cox \(1972\)](#) operationalizes this separation through the heterogeneity in the coefficients $\beta_{K,j}$. Its proof is presented in Appendix C.

Lemma 7 (Structural identifiability and asymptotic learning). Reference framework. Part (i) relies on [Cox \(1972\)](#) and [Bradburn et al. \(2003\)](#) because the daily record Y_t is interpreted as the output of a proportional competing risks system: the heterogeneity in $\beta_{K,j}(\theta)$ alters intensities, hazards, and masses on the same support, allowing types to be contrasted in total variation period by period. Part (ii) resorts to [Kakutani \(1948\)](#) and [Shiryaev \(2019\)](#) because, under conditional independence of the daily experiments and recurrent separation in total variation, the product of affinities vanishes, and the trajectory laws become mutually singular. Finally, [Schervish \(1995\)](#) fixes the language of Bayesian identification under which this singularity translates into the exclusion of false hypotheses and the concentration of the posterior on the true type.

Antecedents and necessary conditions (domain of contrast). The lemma relies on the finite menu Θ_K , the daily record $Y_t = (m_t, d_t)$, the public filtration $\{\mathcal{F}_t^{\text{pub}}\}$, and the predictive kernel of the mechanism, already formalized in Section 2 and in Definitions 1 and 3. In this Lemma, $p_\theta(\cdot \mid \mathcal{F}_t^{\text{pub}})$ denotes the mass function of Y_t at the opening of t under the counterfactual type θ , induced by the fixed PBE.

(N1) Finite menu of types. Without $|\Theta_K| < \infty$, the learning cannot be parame-

terized as a finite mixture, nor can counterfactuals be compared type-by-type on the simplex $\Delta(\Theta_K)$.

(N2) Well-defined daily record. Without the PBE inducing $p_\theta(\cdot \mid \mathcal{F}_t^{\text{pub}})$, there is no daily experiment on which to measure total variation.

(N3) Coherence with the mechanism. The laws p_θ must correspond to the daily draw and to the intensities/hazards of Section 2; otherwise, equation (52) would cease to be the operational link with $\beta_{K,j}$.

(i) Let $\theta, \theta' \in \Theta_K$ be distinct. If in a period t the observable executed triple $\tilde{A}_t = (\tilde{a}_t^F, \tilde{a}_t^K, \tilde{a}_t^S)$ ((4)), the optimal instruments (α_t^*, γ_t^*) , and the other co-variates entering $\lambda_j(t)$ coincide under the counterfactual of replacing only the captor's type with θ or θ' , and furthermore there exists $j \in \{1, 2, 3\}$ such that $\beta_{K,j}(\theta) \neq \beta_{K,j}(\theta')$, then the intensities $\lambda_j(t)$ differ between θ and θ' by the factor $\exp(\beta_{K,j}(\theta) - \beta_{K,j}(\theta')) \neq 1$; consequently, the aggregate closure $q(t)$, the causal mix $\xi_j(t)$, or the marginal hazards $\bar{h}_j(t) = q(t)\xi_j(t)$ are altered, and, barring exact degeneracies that equate all masses of Y_t , the distribution of m_t changes in total variation. More precisely, there exists $\varepsilon = \varepsilon(\theta, \theta', t) > 0$ such that

$$\|p_\theta(\cdot \mid \mathcal{F}_t^{\text{pub}}) - p_{\theta'}(\cdot \mid \mathcal{F}_t^{\text{pub}})\|_{\text{TV}} \geq \varepsilon, \quad (52)$$

where $\|\cdot\|_{\text{TV}}$ is the total variation distance between the two laws on Y_t .

(C1) Local contrapositive. If for all $j \in \{1, 2, 3\}$ it were the case that $\beta_{K,j}(\theta) = \beta_{K,j}(\theta')$ in the fixed counterfactual, the multiplicative channel $\exp(\beta_{K,j}(\theta) - \beta_{K,j}(\theta'))$ does not discriminate; part (i) does not invoke type differences in this margin.

(C2) Knife-edges. Even with $\beta_{K,j}(\theta) \neq \beta_{K,j}(\theta')$, if exact compensations equate

all masses of Y_t , the total variation can vanish; these are parametric arrangements excluded from the operational statement.

(ii) Fitted with a public policy path $(\chi_u)_{u \geq 0}$ (the state's reduced form), suppose that, conditional on this path and on $\mathcal{T}$, the observations $\{Y_t\}_{t \geq 0}$ are independent under $\mathbb{P}_{C,\theta}$ for each θ (equivalently, each Y_t follows a law $p_\theta^{(t)}$ that depends on t only via the already fixed policy). For each pair $\theta \neq \theta'$, let $\mathcal{S}_{\theta,\theta'} \subseteq \mathbb{N}$ be the set of periods t in which (52) holds with $p_\theta^{(t)}$ and $p_{\theta'}^{(t)}$. If, in the event $\mathcal{T}$ (Definition 5), the set $\mathcal{S}_{\theta,\theta_K^*}$ is infinite $\mathbb{P}_{C,\theta_K^*}$ -a.s. and there exists $\varepsilon_\theta > 0$ such that $\|p_\theta^{(t)} - p_{\theta_K^*}^{(t)}\|_{TV} \geq \varepsilon_\theta$ for all $t \in \mathcal{S}_{\theta,\theta_K^*}$, then the laws of $(Y_t)_{t \geq 0}$ conditioned on $\mathcal{T}$ and on the policy under $\mathbb{P}_{C,\theta}$ and under $\mathbb{P}_{C,\theta_K^*}$ are mutually singular.

(C3) Infiniteness and margin. Without an infinite subset of periods with $TV \geq \varepsilon_\theta > 0$, the product of Hellinger affinities does not vanish, and the criterion of Kakutani (1948) does not, by itself, force singularity between the product laws.

(C4) Independence conditional on policy. This is a transparent regime sufficiency; outside of it, the law of (Y_t) may not factor as $\bigotimes_t p^{(t)}$, and an alternative route is required (e.g., recurrence of public states where (i) is reactivated with the same bound).

□

With the three ingredients in place, the bounded martingale of Lemma 4, the asymptotic identification of Definition 6, and the structural separation of Lemma 7, the following theorem establishes that the state's beliefs converge almost surely to the true type. Its proof is presented in Appendix C.

Theorem 8 (Identification and Consistency of True Types). Consider a prior $\mu_0 \in \Delta(\Theta_K)$ with full support on the type space and a PBE-consistent Bayesian updating according to Definition 3. Let $\theta^* := \theta_K^*$ be the realized type and $\mu_C := \mathbb{P}_{C, \theta^*}(\cdot \mid \mathcal{T})$ the law of the true type restricted to the prolonged captivity event $\mathcal{T}$ (Definition 5) and renormalized.

If the mechanism is asymptotically identifiable in the event $\mathcal{T}$ under Definition 6, and since the bounded martingale structure of beliefs guarantees convergence to a random limit by Lemma 4, the following properties are satisfied μ_C -almost surely:

1. Exclusion of false hypotheses. The posterior mass assigned to any false type vanishes asymptotically:

$$\mu_t(\theta) \xrightarrow{t \rightarrow \infty} 0, \quad \forall \theta \in \Theta_K, \theta \neq \theta^*.$$

This property is met under the asymptotic sufficiency of Lemma 7, part (ii), where the recurrent separation in total variation and the singularity criterion of Kakutani (1948) ensure that the trajectory laws of different types are mutually singular Shiryaev (2019).

2. Concentration on the true type. The posterior mass of the realized type converges to unity:

$$\mu_t(\theta^*) \xrightarrow{t \rightarrow \infty} 1.$$

This collapse is a direct consequence of the normalization of the probability measure on the finite simplex $\Delta(\Theta_K)$ after the exclusion of false types.

Consequently, on the event $\mathcal{T}$, the sequence of beliefs $\{\mu_t\}$ converges weakly to the Dirac measure δ_{θ^*} .

□

3. Mechanism implementability

3.1. Implementability

The implementability of the mechanism connects two distinct levels. The first is strategic consistency: the existence of a profile in which the state, the captor, and the family are sequentially rational and update beliefs via Bayes' rule at all reached nodes. The second is the sustainability of the public outcome: that the state's social choice rule, along with the instruments (α_t, γ_t) , simultaneously satisfies IC^K , IR^K , and IR^F . Therefore, the following definitions and Theorem 11 are crucial: they guarantee that the policy subsequently calibrated is not only informative for type identification but also executable as an equilibrium path when the captor's type is private and the family's participation cannot be mechanically imposed.

Definition 9 (Assessment and PBE). Consider a discrete-time Bayesian game with horizon $\mathbb{T} = \{0, 1, \dots, \tau\}$, players $\mathcal{J} = \{S, K, F\}$, and captor's type $\theta_K \in \Theta_K$ according to a common prior $\mu_0 \in \Delta(\Theta_K)$ with $\mu_0(\theta) > 0$ for all $\theta \in \Theta_K$. Following Fudenberg and Tirole (1991a) and Kreps and Wilson (1982), an assessment in this dynamic game is described by the following formal objects:

- *Strategy profile*: $s = (s^S, s^K, s^F)$, where $s^S = \chi = (\chi_t)_{t \in \mathbb{T}}$ with

$$\chi_t: H_t \times \mathcal{R} \longrightarrow \mathcal{A}^{S, \text{disc}} \times \mathcal{A}^\alpha \times \mathcal{A}^\gamma, \quad \mathcal{A}^{S, \text{disc}} = \{\text{Rescue}, \text{Negotiate}\}.$$

The profiles s^K and s^F are behavioral strategies of the private agents in their respective information sets.

- *Belief system*: $\mu = \{\mu_t\}_{t \geq 0}$, with $\mu_t \in \Delta(\Theta_K)$, summarizing the uncertainty regarding θ_K after what has been publicly observed up to the opening of t .

- *Ex-ante law*: $\mathbb{P}_E^s$ is the probability law of the episode induced by the profile s , the prior μ_0 , and the physical dynamics of captivity.

Perfect Bayesian Equilibrium. The pair (s, μ) is a PBE if conditions (a) and (b) hold:

- (a) Sequential rationality. In every t and in every information set reached with strictly positive probability under $\mathbb{P}_E^s$, the equilibrium actions of each agent $i \in \mathcal{J}$ solve the optimization program conditioned on $\mathcal{I}_t^i$, given μ_t and s^{-i} .

For the state:

$$(a_t^{S*}, \alpha_t^*, \gamma_t^*) \in \arg \min_{(a_t^S, \alpha_t, \gamma_t) \in \Gamma_t(\mu_t)} \left\{ \sum_{\theta \in \Theta_K} \mu_t(\theta) L_t(a_t^S, \alpha_t, \gamma_t, \theta, \iota_t) + \Pi_{t,a^S}^S(\alpha_t, \gamma_t; \mu_t) - \psi_H \Delta H_t(\alpha_t, \gamma_t) \right\},$$

with the selection rule $a_t^{S*} = \text{Rescue}$ if the expected cost of rescue does not exceed that of negotiation, and $a_t^{S*} = \text{Negotiate}$ otherwise, consistent with the discrete rule (46). For each private agent $i \in \{K, F\}$:

$$a_t^{i*} \in \arg \max_{a_t^i \in \mathcal{A}^i} \mathcal{U}_t^i(a_t^i \mid \mathcal{I}_t^i, \chi, s^{-i}),$$

where L_t is the type-conditional social loss, Π_{t,a^S}^S is the penalty for deviation from the Bayesian benchmark of the chosen branch, and $\mathcal{U}_t^K, \mathcal{U}_t^F$ are the expected utilities of the captor and the family in t .

- (b) On-path Bayesian consistency. In every node reached with positive probability under $\mathbb{P}_E^s$, beliefs coincide with the Bayesian inference from μ_0 and the observed public signals. For all $\theta \in \Theta_K$ and all $t \geq 1$,

$$\mu_t(\theta) = \mathbb{P}_E^s(\theta_K = \theta \mid \mathcal{F}_t^{\text{pub}}), \quad \mathcal{F}_t^{\text{pub}} = \sigma(z, \zeta_0, \dots, \zeta_{t-1}), \quad (53)$$

where $\mathbb{P}_E^s$ is the ex-ante law induced by the strategy profile s , the common prior μ_0 , and the physical dynamics of captivity. The definition is read as a fixed-point condition: given s , the law $\mathbb{P}_E^s$ determines the beliefs by Bayes at reached nodes; given this belief system, s must be sequentially optimal. At $t = 0$, μ_0 is the common prior conditioned on the observable geographic zone. Off the equilibrium path, beliefs are admissible and compatible with sequential rationality in the remaining subgames [Fudenberg and Tirole \(1991a\)](#); [Kreps and Wilson \(1982\)](#); the subsequent result is one of existence, not uniqueness, and therefore admits the standard off-path multiplicity in dynamic games with incomplete information.

□

Definition 10 (Nash-Bayes- IC^K, IR^K, IR^F Implementability). Let (s^*, μ^*) be a PBE in the sense of Definition 9. The mechanism is implementable Nash-Bayes- IC^K, IR^K, IR^F if the equilibrium profile additionally satisfies the following three feasibility conditions.

Bayesian incentive compatibility of the captor (IC^K). For all $t \in \mathbb{T}$ and all rival mimicry $\theta_j \in \Theta_K$,

$$\mathbb{E}_{\theta \sim \mu_t^*} [V_t^K(a^*(\theta) \mid \theta, \alpha_t^*, \gamma_t^*) - V_t^K(a^*(\theta_j) \mid \theta, \alpha_t^*, \gamma_t^*)] \geq 0,$$

where $a^*(\theta)$ denotes the trajectory of equilibrium actions of type θ under policy χ^* . This is the global non-mimicry condition in expected value: the expectation is taken over the true type, and the deviation θ_j remains fixed.

Bayesian participation constraint of the captor (IR^K). This constraint is not a reservation rent in the sense of [Myerson \(1979\)](#); in kidnapping for ransom, the interaction is coercive and not contractual. It requires that the utility of release

dominates in expected value the best criminal alternative between continuation and death, evaluated at the equilibrium instruments:

$$\mathbb{E}_{\theta \sim \mu_t^*} [U_{\text{rel}}^K(\theta) - \max \{V_{\text{cont},t}^K(\theta, \alpha_t^*, \gamma_t^*), U_{\text{kill}}^K(\theta, \theta_S)\}] \geq 0.$$

The reading is consistent with the logic of [Becker \(1968\)](#) adapted to armed conflict: public policy erodes the expected profitability of persisting or escalating violence relative to surrender. Here $V_{\text{cont},t}^K$ is the continuation component of the value function V_t^K used in IC^K , evaluated on the branch where captivity has not been absorbed, according to (34).

Participation constraint of the family (IR^F).

$$\mathcal{U}_t^F(a_{\text{coop}}) \geq \mathcal{U}_t^F(a_{\text{col}}),$$

evaluated with the information $\mathcal{I}_t^F$ and the reference equilibrium (s^*, μ^*) , according to (49). The simultaneous compatibility of this constraint with IC^K and IR^K is not assumed separately. The set of policy triples that simultaneously satisfy the three preceding conditions is called the feasible set of policies and is denoted by

$$\Gamma_t(\mu_t) := \left\{ (a_t^S, \alpha_t, \gamma_t) \in \mathcal{A}^{S, \text{disc}} \times [0, 1] \times [0, 1] : IC^K, IR^K, IR^F \text{ are satisfied} \right\}. \quad (54)$$

Building on the existence framework of [Kadan et al. \(2017\)](#), Theorem 11 treats the non-emptiness of $\Gamma_t(\mu_t)$ as the primitive robust-feasibility condition ensuring the existence of an implementable PBE. The proof is presented in Appendix C.

□

Theorem 11 (Conditional implementability of the mechanism in PBE). Consider the PBE of Definition 9 under the following conditions:

- (i) $|\Theta_K| < \infty$, with strictly positive prior $\mu_0(\theta) > 0$ for all $\theta \in \Theta_K$.
- (ii) The action spaces $\mathcal{A}^i$ are finite for all $i \in \{K, S, F\}$ and the game tree has a finite number of nodes in each period t , so that the behavioral strategy space of each agent is a finite product of simplexes.
- (iii) The implementation law $\mathbb{P}_{I,i}$ of the MDG process assigns strictly positive probability to every action $a \in \mathcal{A}^i$ for every plan a_t^{i*} and state X_t (full support, (4)), which corresponds to the base specification with $\underline{c} > 0$.
- (iv) The instrument sets are $\mathcal{A}^\alpha = [0, 1]$ and $\mathcal{A}^\gamma = [0, 1]$, compact and convex; the functions L_t and $\mathcal{U}_t^i$ are measurable, uniformly bounded in $\mathbb{T}$, and continuous in the continuous decision variables. The physical outcome rule G_t is measurable, and the transition kernel Q_t induced by G_t and by the competing-risk probabilities is continuous in those variables.
- (v) Robust feasibility. For each t and each $\mu_t \in \Delta(\Theta_K)$, the feasible set $\Gamma_t(\mu_t)$ defined in (54) is non-empty and compact. For each discrete branch a_t^S , its continuous section

$$\Gamma_t^{a_t^S}(\mu_t) := \{(\alpha_t, \gamma_t) : (a_t^S, \alpha_t, \gamma_t) \in \Gamma_t(\mu_t)\}$$

is convex, and the correspondence $\mu_t \mapsto \Gamma_t(\mu_t)$ is continuous (upper and lower hemicontinuous) and compact-valued in $\Delta(\Theta_K)$. Non-emptiness is a primitive condition of simultaneous compatibility among IC^K , IR^K , and IR^F (Kadan et al., 2017).

- (vi) Parametric separability. For all $t \in \mathbb{T}$, the captor's value function admits

the representation

$$V_t^K(a^*(\theta') \mid \theta_K, \alpha_t, \gamma_t) = \sum_j w_j(a^*(\theta'), \alpha_t, \gamma_t) \cdot \exp(\beta_{K,j}(\theta_K) \cdot z),$$

where the weights $w_j(\cdot)$ depend on the action path $a^*(\theta')$ and the instruments, but not on θ_K . This condition allows auditing the economic direction of the incentives; it does not replace the verification of IC^K in (v) nor the statistical identification of Lemma 7.

- (vii) Quasiconvexity of the restricted objective. For each t , each $\mu_t \in \Delta(\Theta_K)$, and each discrete branch $a_t^S \in \mathcal{A}^{S,\text{disc}}$, the functional

$$(\alpha_t, \gamma_t) \mapsto \sum_{\theta_K \in \Theta_K} \mu_t(\theta_K) L_t(a_t^S, \alpha_t, \gamma_t, \theta_K, \iota_t) + \Pi_{t,a^S}^S(\alpha_t, \gamma_t; \mu_t) - \psi_H \Delta H_t(\alpha_t, \gamma_t)$$

is quasiconvex in (α_t, γ_t) over $\Gamma_t^{a^S}(\mu_t)$. Combined with the convexity of $\Gamma_t^{a^S}(\mu_t)$ (hypothesis (v)), this condition guarantees that the optimal restricted correspondence $\Psi_t(\mu_t) := \arg \min_{\Gamma_t^{a^S}(\mu_t)} \{\cdot\}$ is convex-valued, which makes Kakutani's fixed-point theorem applicable to the state's program. It is satisfied in particular when the complete objective, including the informational motive $-\psi_H \Delta H_t$, preserves quasiconvexity in (α_t, γ_t) ; for example, when L_t and Π_{t,a^S}^S are convex, as is the case with the type-specific quadratic functionals C_{ops} and C_{maint} of the model and with the quadratic deviation penalty, and the entropy term does not destroy the geometry of the relevant sublevel sets.

Then there exists a pair (s^*, μ^*) , with $s^* = (\chi^*, s^{K*}, s^{F*})$, that:

- (a) Constitutes a PBE in the sense of Definition 9: the strategies in s^* are mutual best responses given μ^* , and the beliefs μ_t^* are consistent with Bayes' rule at reached nodes with positive probability.

- (b) Implements the social choice rule f_t : on the equilibrium path, $\chi_t^*(h_t) = (a_t^{S^*}, \alpha_t^*, \gamma_t^*)$ coincides with the allocation that minimizes the expected social loss (45) subject to $\Gamma_t(\mu_t^*)$, in accordance with rule (46).
- (c) Satisfies IC^K , IR^K , and IR^F under the profile (s^*, μ^*) .

□

Theorem 11 bridges learning and strategy by complementing identification. The Kakutani route establishes that sustained pressure can make daily signals discriminate types; this theorem supplies the strategic counterpart, showing that the state can deploy such pressure on an equilibrium path without violating IC^K , IR^K , or IR^F , so the captor does not rationally exit through premature lethal closure and the family does not abandon the mechanism. The guarantee is conditional on robust feasibility, hypothesis (v), verified in the calibrated scenarios though not assured a priori along every belief trajectory: while $\Gamma_t(\mu_t)$ remains non-empty and negotiation incentives dominate the lethal outside option, an implementable PBE keeps adversary conduct predictable enough for evidence on θ_K^* to accumulate. The calibration that follows therefore evaluates a trajectory that is simultaneously strategic and informative.

3.2. Economic policy implications

The implications for economic policy concentrate on how the state should adjust operational pressure when learning about the captor's type. The pressure γ_t does not have a mechanical, uniform effect on the duration of captivity: its impact depends on whether it accelerates the closing channels more than it shifts or inhibits other outcomes. Furthermore, when the posterior assigns greater mass to high-risk types, optimal policy must respond more intensely only if the net marginal

effect of the full program favors greater pressure. The following two propositions formalize these margins: the first links Bayesian identification and optimal pressure as a local result of the state's complete objective; the second characterizes how pressure alters the expected duration of captivity.

The first proposition converts Bayesian learning into a local operational rule: if the public history shifts the posterior toward a type with greater propensity for critical outcomes, state pressure increases when material loss, the deviation premium, and informational value jointly point in that direction. The dynamic acquisition of information of [Zhong \(2022\)](#) provides the foundation for treating information as an intertemporal choice object: the decision maker may accept current learning costs when present instruments modify the quality of future signals and, therefore, the value of later decisions. In this model, that intuition is incorporated through $-\psi_H \Delta H_t$ and the net monotonicity condition (56).

Proposition 12 (Optimal operational pressure and Bayesian type identification). *Suppose $\Theta_K = \{\theta_{low}, \theta_{high}\}$ and fix a discrete branch $b \in \{R, N\}$ in a neighborhood of the posterior belief considered. Let $x_t = (\alpha_t, \gamma_t)$ and define the complete objective of the branch*

$$J_t^b(x_t, \mu_t) := \sum_{\theta \in \Theta_K} \mu_t(\theta) L_t(b, \alpha_t, \gamma_t, \theta, \iota_t) + \Pi_{t,b}^S(\alpha_t, \gamma_t; \mu_t) - \psi_H \Delta H_t(\alpha_t, \gamma_t). \quad (55)$$

If (α_t^, γ_t^*) is an interior optimum of $J_t^b(\cdot, \mu_t)$, the constraints IC^K , IR^K , and IR^F do not locally change regime, the Hessian $H_t^b := \nabla_{xx}^2 J_t^b(x_t^*, \mu_t)$ is positive definite, and the net marginal effect of increasing $\mu_t(\theta_{high})$ on the first-order condition for γ_t favors greater pressure, that is,*

$$e'_\gamma (H_t^b)^{-1} \nabla_{x\mu}^2 J_t^b(x_t^*, \mu_t) < 0, \quad e_\gamma = (0, 1)', \quad (56)$$

then the optimal operational pressure of that branch is locally increasing in the posterior belief over the high-risk type:

$$\frac{d\gamma_t^*}{d\mu_t(\theta_{high})} > 0.$$

The structural separation $\beta_{K,j}(\theta_{high}) - \beta_{K,j}(\theta_{low})$ for $j \in \{2, 3\}$ contributes to condition (56) through the ratio

$$\frac{\tilde{\lambda}_j(t \mid \theta_{high})}{\tilde{\lambda}_j(t \mid \theta_{low})} = \exp(\beta_{K,j}(\theta_{high}) - \beta_{K,j}(\theta_{low})), \quad j \in \{2, 3\},$$

but the monotonicity of γ_t^* depends on the joint net effect of material losses, deviation penalty, and informational gain.

□

The second proposition translates that pressure rule into a directly interpretable magnitude for public policy: the expected duration of captivity. Its importance is that it separates the cases in which intensifying the operation accelerates closure from those in which it may prolong negotiation by displacing the payment channel without sufficiently compensating for the risks of rescue or death. In this sense, the proposition provides the physical translation of the dynamic information design of [Zhong \(2022\)](#): while there the optimal acquisition of information depends on the frequency, precision, and temporal cost of signals, here the net sensitivity κ_h shows how operational pressure alters the probabilistic frequency of outcomes. State intervention can accelerate type revelation by modifying closing hazards, but may also accept the risk of prolonging captivity when that temporal cost allows the structural fingerprint of the captor to be more clearly isolated.

Proposition 13 (Expected duration of captivity and operational pressure). *Standing assumptions. Two working assumptions are imposed. (S1) The derivative*

$\partial \mathbb{E}[\tau \mid \theta_K] / \partial \gamma_t^*$ is a partial derivative: γ_t^* is perturbed in a single period t while the instrument sequence $\{\gamma_u^*\}_{u \neq t}$ is held fixed; expected duration is expressed as $\mathbb{E}[\tau \mid \theta_K] = \sum_{u=0}^{\infty} S(u \mid \theta_K)$ with $S(u \mid \theta_K) = \prod_{r=1}^u p_{\text{Cont},r|\theta_K}$, so only factors $S(u)$ with $u \geq t$ depend on γ_t^* . (S2) Operational pressure shifts each intensity proportionally to its level: $\partial \tilde{\lambda}_j(t \mid \theta_K) / \partial \gamma_t^* = \zeta_{\gamma,j} \tilde{\lambda}_j(t \mid \theta_K)$, where $\zeta_{\gamma,j} := \partial \ln \tilde{\lambda}_j(t \mid \theta_K) / \partial \gamma_t^*$ is the semi-elasticity of intensity j with respect to γ_t^* ; λ_4 is invariant to γ_t^* because it is exogenous.

Define the net sensitivity of type θ_K to the optimal operational pressure γ_t^* as:

$$\kappa_h(\theta_K, t) := \left(\zeta_{\gamma,2} \tilde{\lambda}_2(t \mid \theta_K) + \zeta_{\gamma,3} \tilde{\lambda}_3(t \mid \theta_K) \right) - \zeta_{\gamma,1} \tilde{\lambda}_1(t \mid \theta_K). \quad (57)$$

If $\kappa_h(\theta_K, t) > 0$ for all t , the expected duration of captivity $\mathbb{E}[\tau \mid \theta_K]$ is strictly decreasing in γ_t^* . If $\kappa_h(\theta_K, t) < 0$ for all t , it is strictly increasing. The model thus generates the following testable differential prediction:

$$\text{sgn} \left(\frac{\partial \mathbb{E}[\tau \mid \theta_K]}{\partial \gamma_t^*} \right) = -\text{sgn}(\kappa_h(\theta_K, t)), \quad (58)$$

with the sign varying across types through the structural parameters $\beta_{K,j}(\theta_K)$.

□

4. Application: Kidnapping for Ransom in Colombia

The preceding sections state the mechanism for a general unknown type. This section specializes the primitives to kidnapping for ransom in Colombia and documents the empirical mapping from data to structural parameters. The goal is not to replace the theorems with a case study, but to show that the abstract objects—type menu, priors, hazards, instruments, and feasible set—are recoverable from observable institutional variation.

4.1. Empirical anchoring of the parameter space

Between 1964 and 2025, Colombia registered 61,341 kidnapping cases (CNMH, 2026), concentrated in specific municipalities and closely tied to the evolution of the internal armed conflict; the 1996 creation of the Unified Action Groups for Personal Liberty (GAULA, CRC, Law 282 1996) marked the state’s institutional response to this dynamic, and the resulting variation in coercive capacity, legal reform, and organizational structure across armed groups is precisely the source of the type heterogeneity the mechanism is designed to screen. This applied motivation is developed at length, together with the underlying macroeconometric analysis (a Vector Error Correction model and a Local Linear Trend model of kidnapping’s stochastic trend, formally reconciled by the trend-equivalence corollary), in Appendix A; the full microeconometric estimation underlying the table below is presented in Appendix B.

The parameters that feed the mechanism’s state space and technologies are obtained from a Multinomial Logit (MNL) model of case outcomes and a cause-specific proportional hazards model of Cox (1972), estimated on individual-level microdata from the National Center for Historical Memory (CNMH, 2026) ($n = 1,125$ analytical sample; full specification, robustness checks, and sample-selection diagnostics in Appendix B). Table I consolidates the resulting per-type technology fingerprint for the four captor types $\Theta_K = \{\text{DC}, \text{PAR}, \text{ELN}, \text{FARC}\}$.

Table I: Consolidated per-type hazard and outcome fingerprint (DC, PAR, ELN, FARC)

Type θ_K	Mortality margin (MNL)	Payment hazard ratio (Cox, ref. FARC)	Cum. Rescue @ day 10
FARC	0.002	1.00 (reference)	0.27
ELN	0.007	≈ 1.00 (n.s., $p = 0.317$)	0.28
PAR	0.026	4.28 ($p < 0.001$)	0.33
DC	0.071	1.71 ($p = 0.021$)	n.r. (see Appendix

The mapping from these estimates to the theoretical model of Section 2 is direct. The mortality margins order the reservation utility of release $U_{\text{rel}}^K(\theta) = -\kappa_{\text{rel}}(\theta)$ across types, since organizations with lower logistical retention capacity (higher mortality) face a correspondingly different calculus in the IR^K participation constraint. The Cox hazard ratios are transferred to the mechanism as the type-specific effective intensities $\tilde{\lambda}_j(t \mid \theta_K)$ that govern the daily competing-risks signal ζ_t and, through it, the Bayesian update of μ_t ; their empirical separation across types is the observational counterpart of the uniform total-variation separation hypothesis in Lemma 7, on which Theorem 8 rests. The scenario-specific initial beliefs $\mu_0(\theta \mid z, \theta_V)$ used in the calibration of Section 5 are reported there together with each simulated scenario.

4.2. From estimators to structural inputs

The theoretical model in Section 2 requires two empirical inputs so as not to depend on ad hoc assumptions: a persistent macroeconomic environment that disciplines retention costs and a type structure with differentiated priors, technologies, and hazards. The VEC system, the Gonzalo–Granger decomposition, and the trend-equivalence corollary (Appendix A) provide the first input. The cointe-

gration among K^E , K^{NE} , and r establishes that kidnapping for ransom does not evolve as an isolated series, but as part of a long-run system in which coercive capacity and macro-financial conditions act as permanent forces. The equivalence between the VEC–GG trend and the LLT trend prevents calibration from depending on a single econometric specification, while the adjustment structure of the VEC shows that K^E is the margin that absorbs disequilibria. Therefore, the state vector of the mechanism, the feasible set $\Gamma_t(\mu_t)$, and the constraints IC^K and $IR^{K,F}$ are anchored in persistent regularities, not in a purely theoretical normalization.

On this basis, the MNL and the Cox model connect microeconomic evidence with the initial distribution $\mu_0 \in \Delta(\Theta_K)$ and the technology vectors $\beta_{K,j}(\theta_K)$. The MNL provides the static fingerprint of each type by relating perpetrator, victim, context, and tactics to the conditional distribution of outcomes; thus, the initial belief and terminal payments are set with observed weights. The Cox model provides the dynamic dimension: its cause-specific hazards are transferred to the mechanism as rates $\lambda_j(\theta_K, t)$ that govern the daily signal ζ_t , update the posterior μ_t , and determine the speed of learning. The empirical separation of those rates across organizations is the observational basis of Lemma 7 and the criterion of Kakutani (1948): if the trajectories induced by the types generate sufficiently distinct laws, the posterior can concentrate on the true type. Together, the empirical block converts macro and micro data into priors, costs, hazards, and constraints directly usable by the dynamic mechanism. The explicit translation protocol between estimators and structural parameters is detailed at the beginning of Section 5.

5. Calibration of the mechanism and simulations

This section calibrates the specialized mechanism of Section 4 in two simulated episodes of finite horizon, holding fixed the general architecture of Section 2. The exercise combines macro and microeconomic priors and diagnostics to study four objects: the evolution of beliefs $\mu_\tau(\theta)$, the signals observed by the state, the response of the instruments (α_t, γ_t) , and compliance with the IR and IC constraints. *Scenario I* corresponds to ELN, with $\theta_K^* = \text{ELN}$, lax state, low-wealth family, public victim, and Caribbean location; *Scenario II* corresponds to PAR, with $\theta_K^* = \text{PAR}$, lax state, high-wealth family, private victim, and East-Jungle location. Both are simulated over a fixed horizon $\bar{T} = 150$ periods (a calibration parameter distinct from the endogenous stopping time τ defined in (3)), and optimal policy is computed from the state's minimization of social cost. In the order (DC, PAR, ELN, FARC), the initial priors are $\mu_0^{\text{ELN}} = (0.25, 0.35, 0.05, 0.35)$ and $\mu_0^{\text{PAR}} = (0.55, 0.05, 0.35, 0.05)$. In the figures, the gray bar marks the realized stopping time τ (the first period with $m_\tau \in \mathcal{M}_{\text{term}}$); the trajectory after that bar is a counterfactual extension under an absorbing state (the mechanism law conditioned on non-absorption) that illustrates learning and identification dynamics independently of the closure moment, without constituting a real episode trajectory.

What the simulations are designed to show.. Relative to a purely myopic policy that ignores $-\psi_H \Delta H_t$, the calibrated rule internalizes the informational subsidy and reallocates (α_t, γ_t) toward instruments that separate likelihoods. The ELN and PAR exercises therefore illustrate three operational implications of the theory: (i) finite-horizon posterior concentration in the direction predicted by Theorem 8; (ii) MDG-enabled full-support updating along the path; and (iii) ex-post compliance

with IC^K , IR^K , and IR^F under the implementability logic of Theorem 11. They do not claim asymptotic identification on $\mathcal{T}$, which remains a theoretical limit statement.




Figure I: Bayesian beliefs of each type $\mu_\tau(\theta)$ (ELN: $\tau \leq 150$; PAR: $\tau \leq 150$).

Note: Panel (a) corresponds to Scenario I: $\theta_K^* = \text{ELN}$, lax θ_S , low-wealth θ_F , public θ_V , and Caribbean z . Panel (b) corresponds to Scenario II: $\theta_K^* = \text{PAR}$, lax θ_S , high-wealth θ_F , private θ_V , and East-Jungle z . The gray band marks τ^{hyp} , the first period in which the mechanism outcome differs from continuation.

The posterior of the true type allows evaluating the capability of the mechanism to separate types from the sequence of signals. In the ELN scenario, reported in Figure I, posterior mass shifts gradually toward ELN until reaching $\mu_{150}(\text{ELN}) \approx 1,000$; this implies that the mechanism accumulates identification even in an environment with lower family wealth and a public victim. In PAR, the posterior of the true type starts from $\mu_0(\text{PAR}) = 0,050$ and reaches $\mu_{150}(\text{PAR}) \approx 1,000$; the trajectory is not monotone in the initial phase, because the event combines early signals compatible with other types and then shifts mass toward PAR. The gray bars mark the realized stopping time ($\tau = 16$ in ELN, $\tau = 5$ in PAR); the subsequent counterfactual extension shows that the speed of

learning and the absorption moment are distinct objects of the model. Taken together, the figure illustrates finite-horizon behavior: posterior mass shifts toward the true type as signals accumulate separation, in the direction predicted by Theorem 8 under $\mathcal{T}$, without constituting direct evidence for the asymptotic result; the speed of convergence depends on the institutional, territorial, and case-visibility environment.




Figure II: Operational pressure γ_t^* and benchmark γ^R against τ .

Note: Panel (a) shows the operational pressure trajectory in Scenario I (ELN). Panel (b) shows the equivalent trajectory in Scenario II (PAR). The dashed line corresponds to the Rescue benchmark γ^R . The gray band marks τ^{hyp} .

The optimal operational pressure allows evaluating how the state adjusts the intensity of intervention as it learns about the captor's type. Figure II compares the trajectory of γ_t^* with the Rescue benchmark γ^R , which takes the value $\gamma^R = 0,92$ in ELN and $\gamma^R = 0,86$ in PAR. In the ELN scenario, γ_t^* is located on average around $\bar{\gamma}^* \approx 0,908$ and remains close, though below, the Rescue benchmark, consistent with a Caribbean environment and lax state. In PAR, pressure remains near the Rescue benchmark and reaches $\bar{\gamma}^* \approx 0,868$, consistent with an East-Jungle environment and lax state where intervention retains value even when the focal outcome of the event is payment. In both cases, γ_t^* does not mechanically replicate

either the Rescue or the Negotiate branch: dual control combines exploitation of the instrument and expected information, so the series fluctuates between benchmarks according to the material risk of the episode and the value of separating types.

The optimal financial interdiction summarizes the intensity with which the state restricts payment capacity and economic coordination during the episode. Figure III compares the trajectory of α_t^* with the Rescue benchmark α^R , which takes the value $\alpha^R = 0,90$ in ELN and $\alpha^R = 0,84$ in PAR. In ELN, $\bar{\alpha}^* \approx 0,908$ and the instrument remains at elevated levels, compatible with a public victim and low-wealth case where the block operates more as strategic containment than as closure of a high payment capacity. In PAR, $\bar{\alpha}^* \approx 0,901$ and the trajectory remains elevated over much of the support, indicating that the block retains high strategic value when the state is lax, the family has high wealth, and the East-Jungle region opens room for payments. Thus, α_t^* does not mechanically replicate a pure branch: it moves between Rescue and Negotiate benchmarks according to accumulated information, family payment capacity, and the value of inducing separation among types.

The posterior precision, defined as $\iota_t = \max_{\theta \in \Theta_K} \mu_t(\theta)$, allows observing how operational pressure changes when the state gains confidence about the identity of the true type. Figure IV reorders the simulation on the plane (ι_t, γ_t^*) ; the arrows indicate the advance of τ and avoid interpreting the observations as a linear interpolation. In ELN, posterior precision increases gradually while operational pressure remains high but contained, reflecting learning without immediate collapse toward a single dominant signal. In PAR, the trajectory concentrates in a zone of high pressure and rising precision, consistent with an episode in which

information requires a longer window before resolving the captor's identity.




Figure III: Financial interdiction α_t^* and benchmark α^R against τ .

Note: Panel (a) presents the optimal financial interdiction in Scenario I (ELN). Panel (b) presents the same instrument in Scenario II (PAR). The dashed line corresponds to the Rescue benchmark α^R . The gray band marks τ^{hyp} .

Financial interdiction against posterior precision allows evaluating whether learning reduces or reinforces the need to restrict payment capacity. In the ELN scenario, reported in Figure V, α_t^* remains high while ι_t increases, indicating that identifying the insurgent type does not eliminate the value of maintaining preventive interdiction in a public-victim, low-wealth case. Something similar occurs in PAR: although the posterior accumulates precision, α_t^* remains at high levels over much of the support. This pattern is consistent with environments where the state is lax and financial interdiction retains relevance as an instrument of control, payment containment, and induction of separation among types.






Figure IV: Trajectory (l_t, γ_t^*) with arrows indicating temporal advance.

Note: Panel (a) corresponds to Scenario I (ELN) and panel (b) to Scenario II (PAR). Each trajectory shows 20 arrows traced over homogeneous intervals of τ ; the circle identifies the end point of the simulation, $\tau = 150$.

The constraints diagnostic confirms that the strategic feasibility of the mechanism is preserved in both scenarios. As Figure VI shows, the captor's incentive constraint (IC^K), the captor's participation constraint (IR^K), and the family's participation constraint (IR^F) are satisfied in all 150 simulated cycles in ELN and PAR. The verification is *a posteriori* with respect to the true type θ_K^* : the optimal instruments $(a_S^*, \alpha_t^*, \gamma_t^*)$ are evaluated directly on the utilities of the true type, without weighting by the belief μ_t . The mechanism is therefore feasible in an ex-post sense: given θ_K^* , optimal policy satisfies incentive compatibility and participation without abandoning the informational component.




(a) Scenario I: ELN




(b) Scenario II: PAR

Figure V: Trajectory (l_t, α_t^*) with arrows indicating temporal advance.

Note: Panel (a) corresponds to Scenario I (ELN) and panel (b) to Scenario II (PAR). Each trajectory shows 20 arrows traced over homogeneous intervals of τ ; the circle identifies the end point of the simulation, $\tau = 150$.



(a) Scenario I: ELN



(b) Scenario II: PAR

Figure VI: Frequency of compliance with IC^K , IR^K , and IR^F (% of cycles; ELN: $\tau \leq 150$; PAR: $\tau \leq 150$).

Note: Panel (a) summarizes compliance with the captor's incentive compatibility, the captor's participation, and the family's participation in Scenario I (ELN). Panel (b) presents the same diagnostic for Scenario II (PAR). Percentages are computed over the simulated cycles of each scenario.

The expected informational gain ΔH_t quantifies the anticipated reduction in entropy induced by the state's instruments, in the sense of the optimal dynamic acquisition of information of [Zhong \(2022\)](#). That informational subsidy concentrates at the start of the simulation, as illustrated in Figure [VII](#): in ELN it reaches its maximum at $\tau = 10$ with $\Delta H_{10} \approx 0,043$, and in PAR at $\tau = 6$ with $\Delta H_6 \approx 0,035$. As ΔH_t converges to zero, additional observations cease to reduce posterior uncertainty and optimal decision comes to be governed by material losses and a belief already concentrated on the true type.





Figure VII: Expected informational gain ΔH_t against τ .

Note: Panel (a) corresponds to Scenario I (ELN) and panel (b) to Scenario II (PAR). The series reports the expected entropy reduction induced by the state's optimal policy in each simulated cycle.

In both scenarios and for all types, $\kappa_h(\theta_K^*, t) > 0$ in 100 % of simulated cycles, as shown in Figure VIII. By the relation $\text{sgn}(\partial \mathbb{E}[\tau \mid \theta_K] / \partial \gamma_t^*) = -\text{sgn}(\kappa_h)$, this implies $\partial \mathbb{E}[\tau \mid \theta_K] / \partial \gamma_t^* < 0$: higher operational pressure *reduces* expected captivity duration without exception. The mechanism does not face a trade-off between exploiting the instrument and learning: operational pressure shortens captivity and, by modifying closure hazards, simultaneously generates signals that separate the likelihoods of rival types and accelerate posterior convergence.



(a) Scenario I: ELN

(b) Scenario II: PAR

Figure VIII: Net sensitivity $\kappa_h(\theta_K^*, t)$ of expected captivity duration with respect to γ_t^* (ELN: $\tau \leq 150$; PAR: $\tau \leq 150$).

Note: Each panel traces $\kappa_h(\theta_K^*, t)$ for the true type of the scenario. A constant value of +1 implies $-\text{sgn}(\kappa_h) = -1$ in 100 % of cycles, that is, $\partial \mathbb{E}[\tau \mid \theta_K] / \partial \gamma_t^* < 0$: operational pressure γ_t^* *reduces* expected captivity duration over the entire support; the horizontal dashed line marks the reference zero.

The hypothetical terminal outcomes differ across scenarios: *Release* in ELN ($\tau^{\text{hyp}} = 16$) and *Death* in PAR ($\tau^{\text{hyp}} = 5$). In both cases the posterior starts from $\mu_0(\theta^*) = 0,050$ and converges to $\mu_{150}(\theta^*) \approx 1,000$, with $\bar{\alpha}^*$ and $\bar{\gamma}^*$ above 0,86 in all scenarios. These results are summarized in Table II.

Table II: Summary of dynamic calibration by scenario

θ_K^*	$\mu_0(\theta^*)$	$\mu_{\tau_{\max}}(\theta^*)$	$\bar{\alpha}^*$	$\bar{\gamma}^*$	τ^{hyp}	$m_{\tau^{\text{hyp}}}$
ELN	0.050	1.000	0.908	0.908	16	Release
PAR	0.050	1.000	0.901	0.868	5	Death

Taken together, the ELN–PAR contrast shows that the same Bayesian mechanism generates distinct learning trajectories, instrument intensities, and constraint viability rates when, jointly, the captor’s true type and the institutional, family, victim, and geographic margins described in the protocol change. The simulations do

not substitute for Theorem 8, which requires recurrent separation in the event $\mathcal{T}$; they illustrate that, with priors and hazards calibrated to Colombian evidence, the entropic subsidy shifts policy away from purely myopic regimes and accelerates or disciplines posterior concentration in finite horizons, in the direction predicted by the identification result.

6. Discussion and conclusions

This paper develops an indirect dynamic mechanism for environments in which equilibrium actions pool and types are not elicited by free reports. Relative to the classical reduction via the Revelation Principle, the contribution is to construct the operational indirect object: public instruments, endogenous implementation noise, and Bayesian learning over action-generated signals.

At the theoretical level, the MDG kernel supplies full-support materialization of latent intentions without an exogenous tremble, keeping Bayes' rule well defined off path. The Type Identification Theorem shows that recurrent total-variation separation of type-contingent observation laws implies mutual singularity of trajectory measures and almost-sure posterior concentration on the true type on the prolonged-interaction event $\mathcal{T}$. The Conditional Implementability Theorem complements identification by guaranteeing a PBE that satisfies incentive compatibility and individual rationality whenever the feasible set remains non-empty. The asymptotic result is intentionally sharp: it clarifies the informational limit of the architecture; finite-horizon behavior is a distinct, applied object.

At the applied level, Colombian kidnapping for ransom maps the abstract type space into estimated criminal technologies. Competing-risks and multinomial evidence discipline hazards and priors; ELN-PAR simulations show that the cali-

brated dual-control rule accelerates posterior concentration, sustains high but disciplined instruments, and preserves IR/IC feasibility, in the direction predicted by the theorems without substituting for them.

Two extensions are natural. The first is a finite-horizon identification bound that translates total-variation separation into explicit posterior concentration rates under endogenous stopping. The second is multi-agent concurrent identification when several privately informed parties interact under a shared public history. Both would transfer the framework beyond adversarial screening to other markets with noisy implementation and observational pooling.

Declaration on the use of Generative Artificial Intelligence

In the preparation of this manuscript, the author used generative artificial intelligence tools, specifically *Cursor* and *Antigravity*, to assist in the following tasks: literature review and text editing; programming in Lean 4 for the formalization and interactive proof of theoretical theorems and lemmas; programming of quantitative and empirical code in Matlab, Stata, and Python; as well as the design and online deployment of the calibrated dynamic model interface via GitHub and Streamlit. Consistent with academic ethical guidelines, the author is solely responsible for the integrity of the document, having thoroughly verified and validated all information, proofs, estimations, and codes provided or generated by these technologies.

References

Abbring, Jaap H. and Gerard J. Van den Berg (2003), “The nonparametric identification of treatment effects in duration models.” *Econometrica*, 71 (5), 1491–

1517.

Aghion, Philippe, Patrick Bolton, Christopher Harris, and Bruno Jullien (1991), “Optimal learning by experimentation.” *The Review of Economic Studies*, 58 (4), 621–654.

Becker, Gary S. (1968), “Crime and punishment: An economic approach.” *Journal of Political Economy*, 76 (2), 169–217.

Bergemann, Dirk and Stephen Morris (2012), “An introduction to robust mechanism design.” *Foundations and Trends in Microeconomics*, 8 (3), 169–230.

Blackwell, David and Lester E. Dubins (1962), “Merging of opinions with increasing information.” *The Annals of Mathematical Statistics*, 33 (3), 882–886.

Bolton, Patrick and Mathias Dewatripont (2005), *Contract Theory*, MIT Press, Cambridge, MA.

Bradburn, MJ, TG Clark, SB Love, and DG Altman (2003), “Survival analysis part ii: Multivariate data analysis – an introduction to concepts and methods.” *British Journal of Cancer*, 89 (3), 431–436.

CNMH (2026), “Portal de datos: El conflicto en cifras.” Conflicto en Cifras.

Courty, Pascal and Hao Li (2000), “Sequential screening.” *Review of Economic Studies*, 67 (4), 697–717.

Cox, D. R. (1972), “Regression models and life-tables.” *Journal of the Royal Statistical Society. Series B (Methodological)*, 34 (2), 187–220.

- CRC, Law 282 (1996), “Law 282 de 1996.” Diario Oficial n.º 42.802. Eva Función Pública.
- Doob, Joseph L. (1953), *Stochastic Processes*, Wiley, New York. Disponible en Internet Archive.
- Easley, David and Nicholas M. Kiefer (1988), “Controlling a stochastic process with unknown parameters.” *Econometrica*, 56 (5), 1045–1064.
- Eső, Péter and Balázs Szentes (2007), “Optimal information disclosure in auctions and the handicap auction.” *Review of Economic Studies*, 74 (3), 705–731.
- Fudenberg, Drew and David K. Levine (1993), “Self-confirming equilibrium.” *Econometrica*, 61 (3), 523–545.
- Fudenberg, Drew and Jean Tirole (1991a), *Game Theory*, MIT Press, Cambridge, MA.
- Fudenberg, Drew and Jean Tirole (1991b), “Perfect bayesian equilibrium and sequential equilibrium.” *Journal of Economic Theory*, 53 (2), 236–260.
- Gibbard, Allan (1973), “Manipulation of voting schemes: A general result.” *Econometrica*, 41 (4), 587–601.
- Holmström, Bengt and Roger B. Myerson (1983), “Efficient and durable decision rules with incomplete information.” *Econometrica*, 51 (6), 1799–1819.
- Jackson, Matthew O. (2001), “A crash course in implementation theory.” *Social Choice and Welfare*, 18 (4), 655–708.

- Jaynes, E. T. (1957), “Information theory and statistical mechanics.” *Physical Review*, 106 (4), 620–630.
- Kadan, Ohad, Philip J. Reny, and Jeroen M. Swinkels (2017), “Existence of optimal mechanisms in principal-agent problems.” *Econometrica*, 85 (3), 769–823.
- Kakutani, Shizuo (1948), “On equivalence of infinite product measures.” *Annals of Mathematics*, 49 (2), 214–224.
- Kalai, Ehud and Ehud Lehrer (1993), “Rational learning leads to nash equilibrium.” *Econometrica*, 61 (5), 1019–1045.
- Kreps, David M. and Robert Wilson (1982), “Sequential equilibria.” *Econometrica*, 50 (4), 863–894.
- Lancaster, Tony (1990), *The Econometric Analysis of Transition Data*, Cambridge University Press, Cambridge.
- Lucas, Robert E. Jr. (1976), “Econometric policy evaluation: A critique.” *Carnegie–Rochester Conference Series on Public Policy*, 1, 19–46.
- Magnac, Thierry and David Thesmar (2002), “Identifying dynamic discrete decision processes.” *Econometrica*, 70 (2), 801–816.
- McCall, John J. (1970), “Economics of information and job search.” *The Quarterly Journal of Economics*, 84 (1), 113–126.
- Myerson, Roger B. (1979), “Incentive compatibility and the bargaining problem.” *Econometrica*, 47 (1), 61–73.

- Myerson, Roger B. (1981), “Optimal auction design.” *Mathematics of Operations Research*, 6 (1), 58–73.
- Myerson, Roger B. and Mark A. Satterthwaite (1983), “Efficient mechanisms for bilateral trading.” *Journal of Economic Theory*, 29 (2), 265–281.
- Pavan, Alessandro, Ilya Segal, and Juuso Toikka (2014), “Dynamic mechanism design: A myersonian approach.” *Econometrica*, 82 (2), 601–653.
- Rochet, Jean-Charles (1987), “A necessary and sufficient condition for rationalizability in a quasi-linear context.” *Journal of Mathematical Economics*, 16 (2), 191–200.
- Rochet, Jean-Charles and Philippe Choné (1998), “Ironing, sweeping, and multi-dimensional screening.” *Econometrica*, 66 (4), 783–826.
- Rust, John (1987), “Optimal replacement of GMC bus engines: An empirical model of Harold Zurcher.” *Econometrica*, 55 (5), 999–1033.
- Schervish, Mark J. (1995), *Theory of Statistics*, Springer Series in Statistics, Springer, New York.
- Selten, Reinhard (1975), “Reexamination of the perfectness concept for equilibrium points in extensive games.” *International Journal of Game Theory*, 4 (1), 25–55.
- Shannon, Claude E. (1948), “A mathematical theory of communication.” *Bell System Technical Journal*, 27 (3), 379–423.
- Shiryaev, Albert N. (2019), *Probability*, 2 edition, Graduate Texts in Mathematics, Vol. 95, Springer, New York.

Sueyoshi, Glenn T. (1992), “Semiparametric proportional hazards estimation of competing risks models with time-varying covariates.” *Journal of Econometrics*, 51 (1-2), 25–58.

Wooldridge, Jeffrey M. (2010), *Econometric Analysis of Cross Section and Panel Data*, second edition, The MIT Press, Cambridge, Massachusetts and London, England.

Zhong, Weijie (2022), “Optimal dynamic information acquisition.” *Econometrica*, 90 (4), 1537–1582.